

Arithmetic for Computers

Lecture 3

September 20th, 2023

Jae W. Lee (jaewlee@snu.ac.kr)

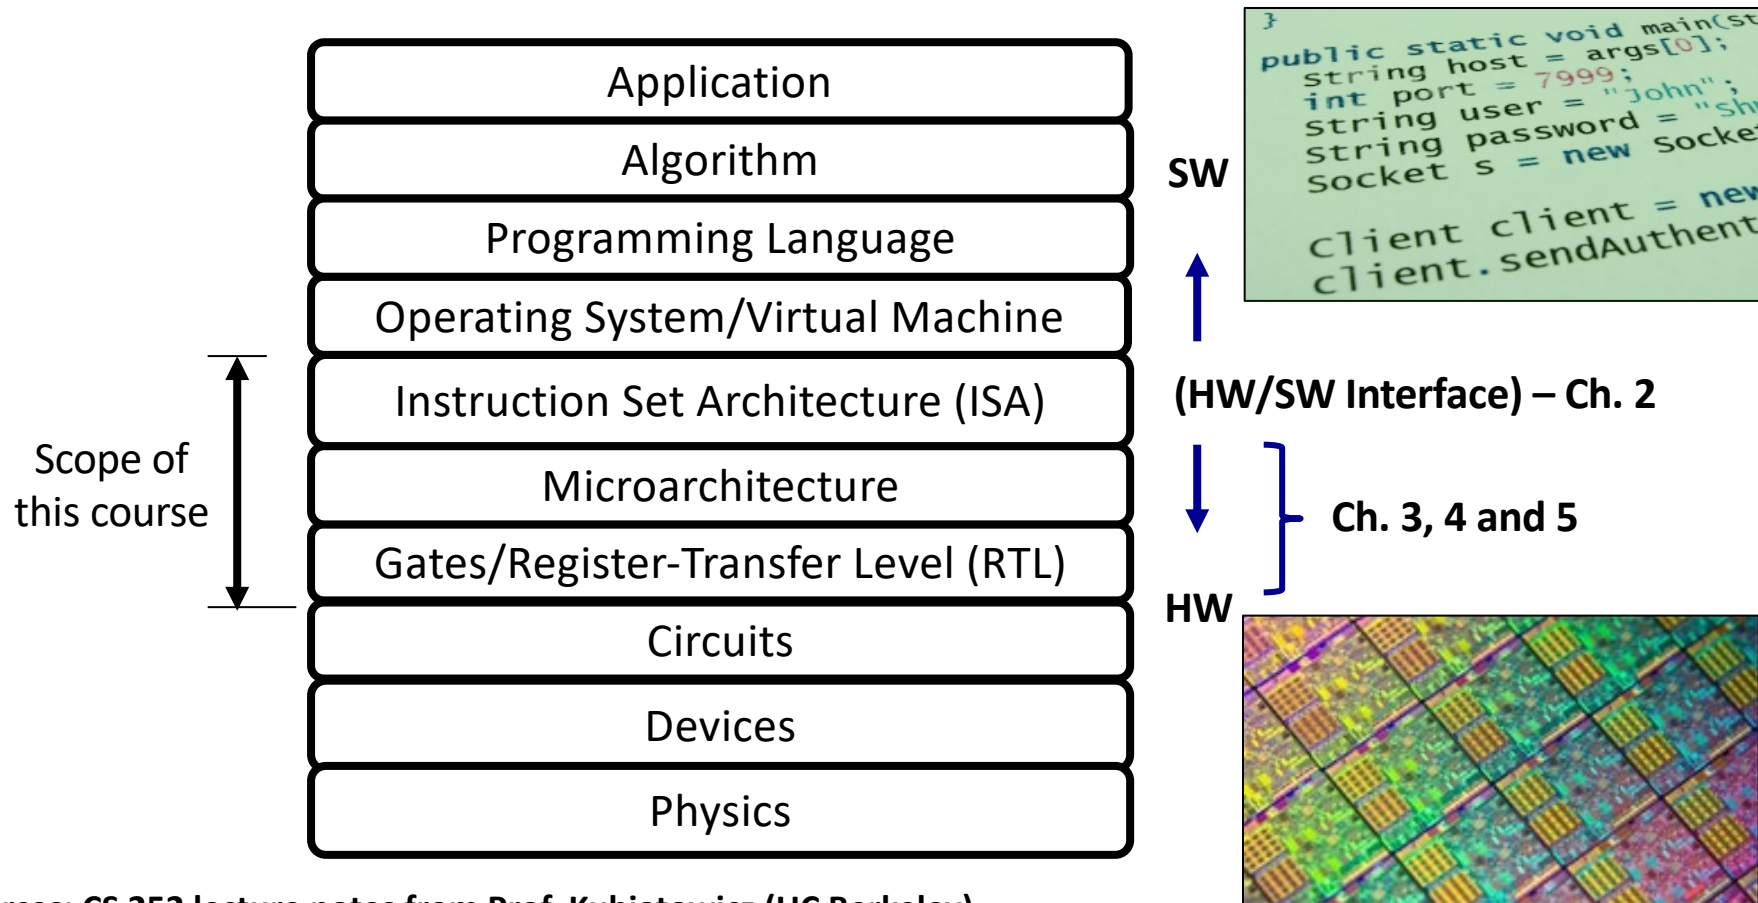
Computer Science and Engineering

Seoul National University

Slide credits: [CS:APP3e] slides from CMU; [COD:RV2e] slides from Elsevier Inc.

Where Are We?

■ Abstraction Layers in Modern Systems



* Sources: CS 252 lecture notes from Prof. Kubiawicz (UC Berkeley).

Coursera lecture notes for HW/SW Interface from Profs. Borriello and Ceze (Univ. of Washington).

Introduction

■ Bits are just bits

- No inherent meaning: conventions define relationship between bits and numbers
- n -bit binary numbers: $0 \sim 2^n - 1$ decimal numbers

■ Of course it gets more complicated

- Numbers are finite (overflow)
- Negative numbers
- Fractions and real numbers
- Precision and accuracy
- Error propagation, ...

Arithmetic for Computers: An Overview

■ Operations on integers

- Textbook: COD 3.1-3.4
- Addition and subtraction
- Multiplication and division
- Dealing with overflow

■ Floating-point real numbers


- Textbook: P&H 3.5
- Representation and operations

■ Subword parallelism

- Textbook: P&H 3.6

Integer Addition

■ Example: 7 + 6



■ Overflow if result out of range

- Adding +ve and -ve operands, no overflow
- Adding two +ve operands
 - Overflow if result sign is 1
- Adding two -ve operands
 - Overflow if result sign is 0

Integer Subtraction

- Add negation of second operand

- Example: $7 - 6 = 7 + (-6)$

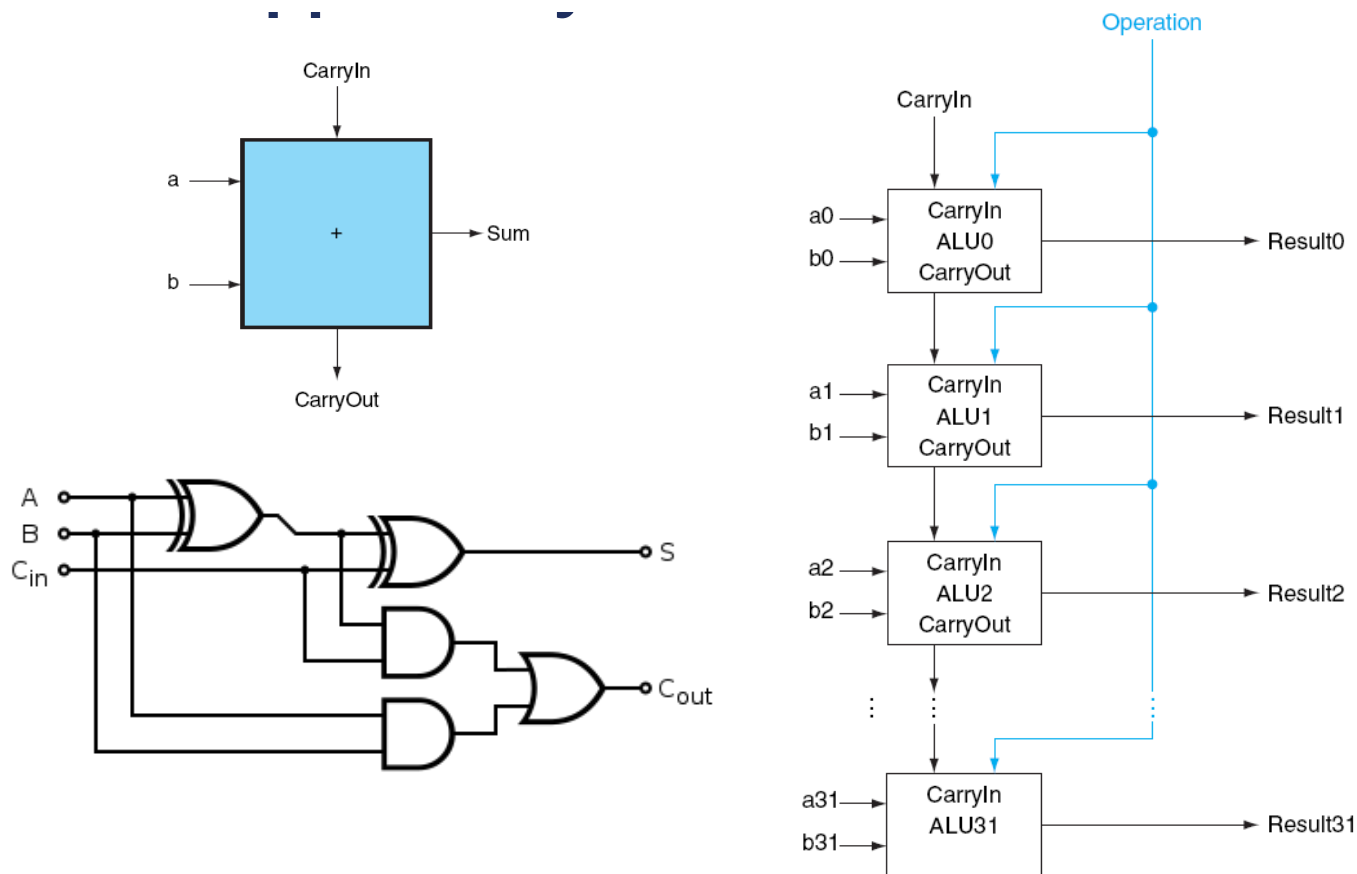
- +7: 0000 0000 ... 0000 0111
- 6: 1111 1111 ... 1111 1010
- +1: 0000 0000 ... 0000 0001

- Overflow if result out of range

- Subtracting two +ve or two -ve operands, no overflow
- Subtracting +ve from -ve operand
 - Overflow if result sign is 0
- Subtracting -ve from +ve operand
 - Overflow if result sign is 1



Simple Adder

■ n -bit ripple-carry adder



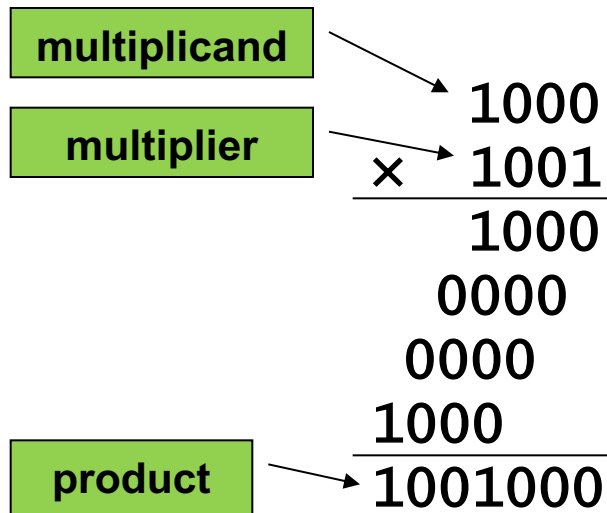
Arithmetic for Multimedia

- **Graphics and media processing operates on vectors of 8-bit and 16-bit data**
 - Use 64-bit adder, with partitioned carry chain
 - Operate on 8×8-bit, 4×16-bit, or 2×32-bit vectors
 - SIMD (single-instruction, multiple-data)
- **Saturating operations**
 - On overflow, result is largest representable value
 - cf. 2s-complement modulo arithmetic
 - e.g., clipping in audio, saturation in video

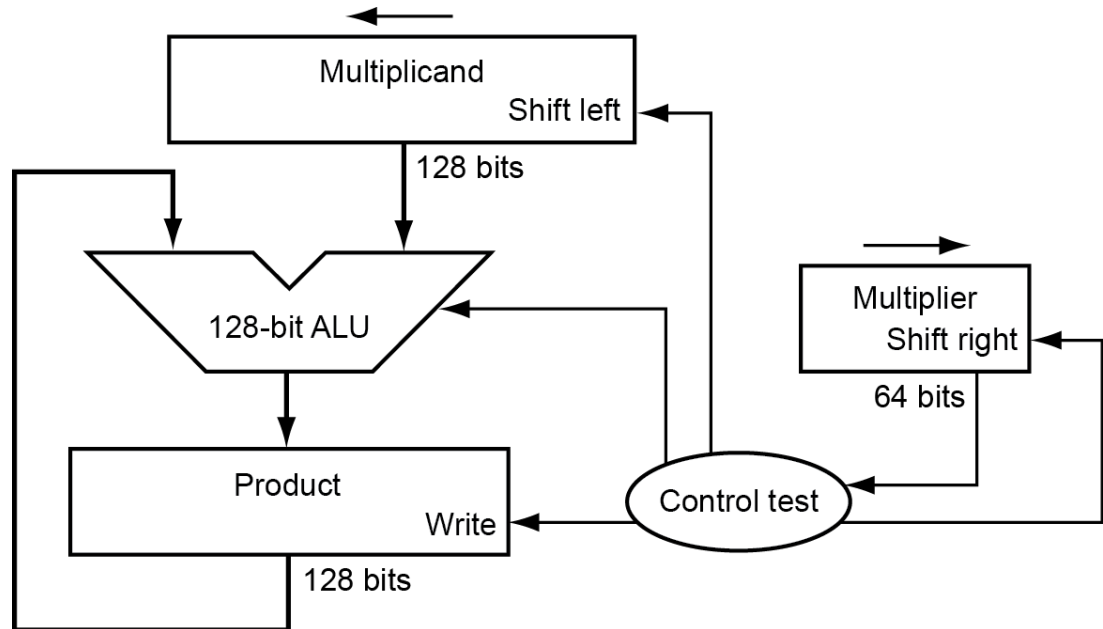


Multiplication

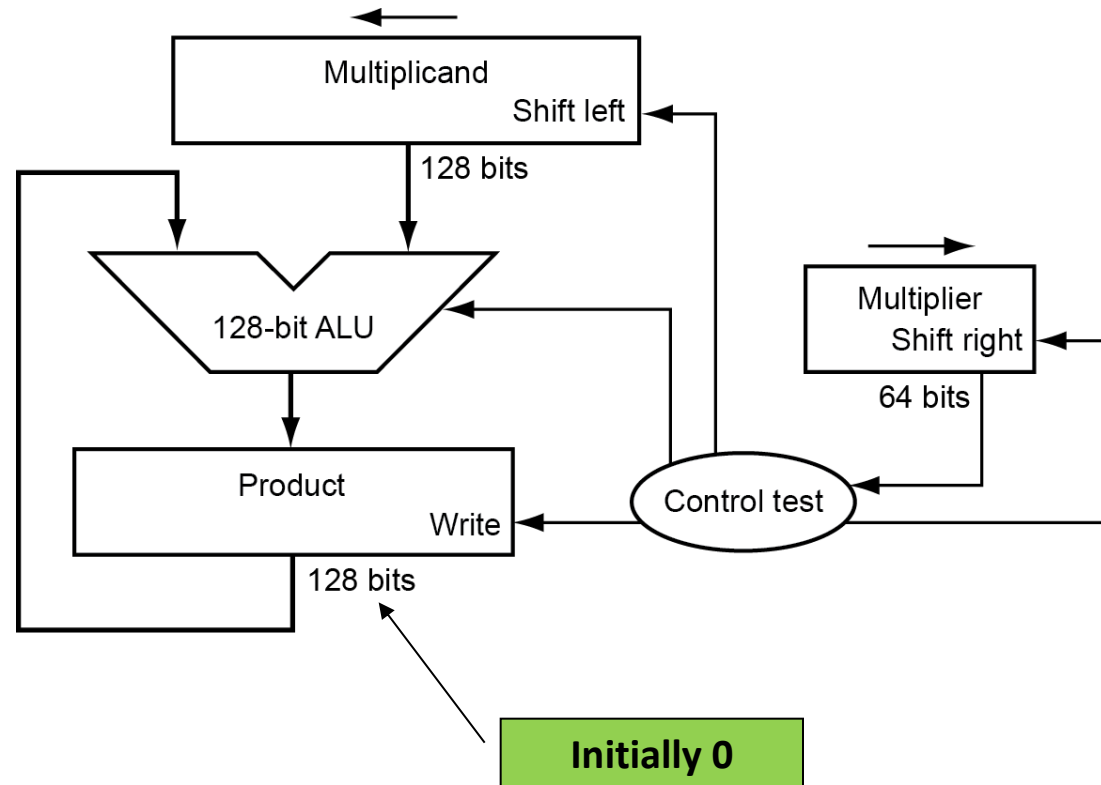
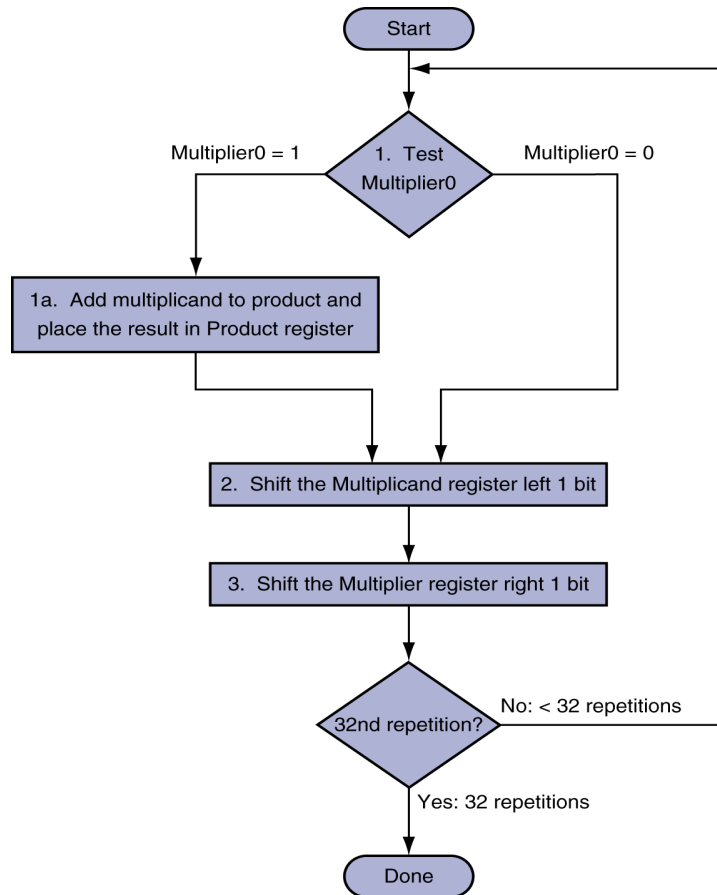
■ Start with long-multiplication approach



Length of product
is the sum of
operand lengths

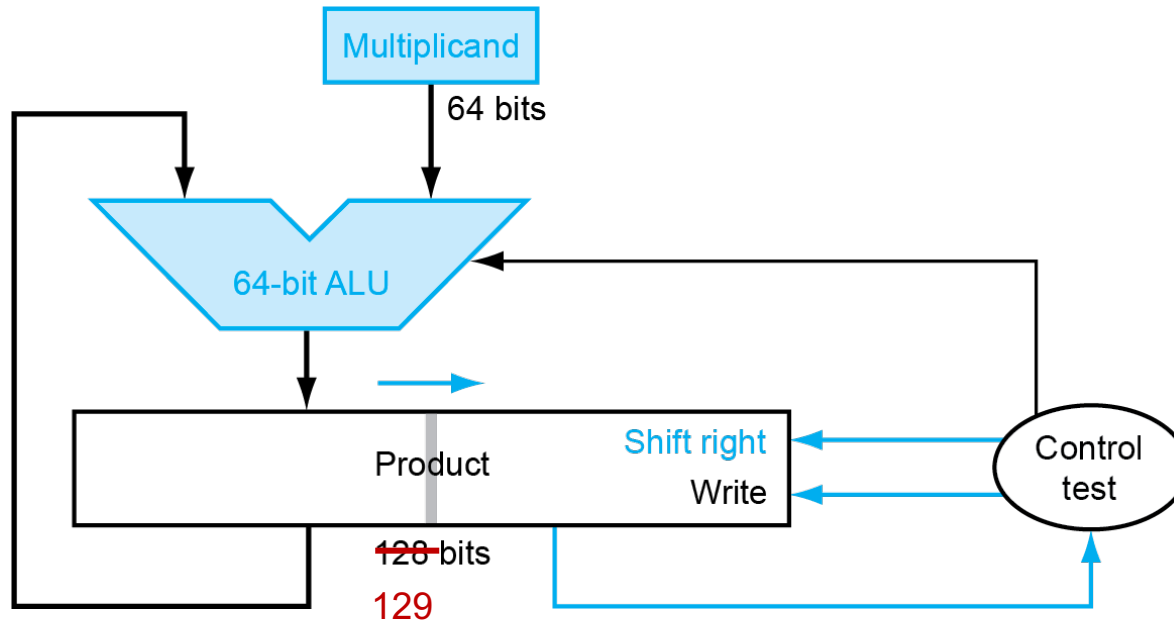


Multiplication Hardware



Optimized Multiplier

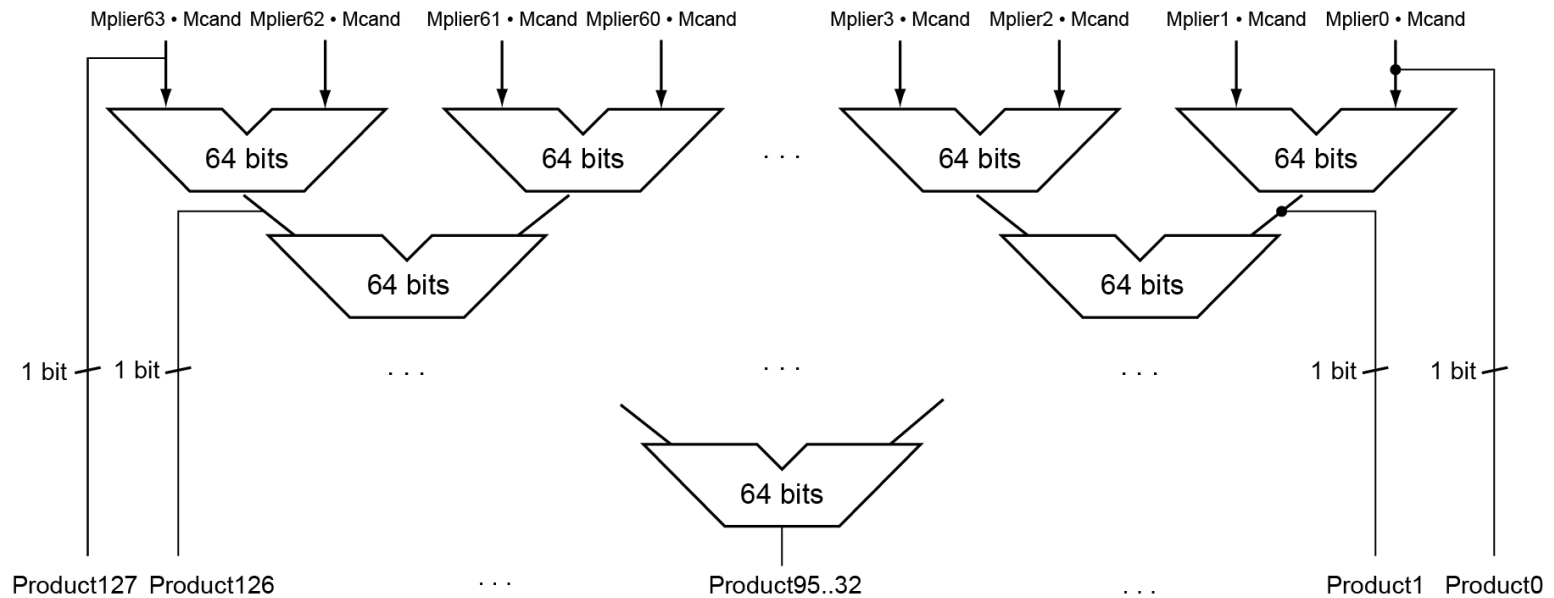
- Perform steps in parallel: add/shift



- One cycle per partial-product addition
 - That's ok, if frequency of multiplications is low

Faster Multiplier

- Uses multiple adders
 - Cost/performance tradeoff



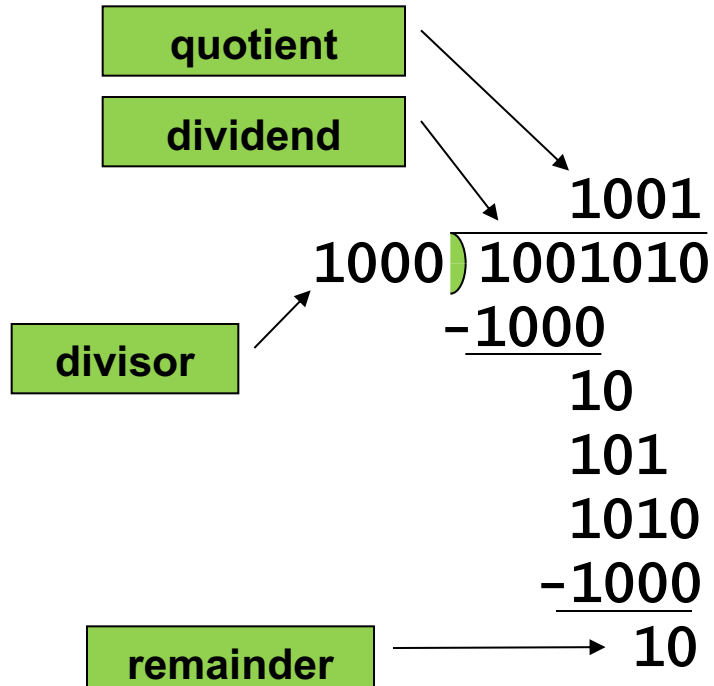
- Can be pipelined
 - Several multiplication performed in parallel

RISC-V Multiplication

■ Four multiply instructions:

- `mul`: multiply
 - Gives the lower 64 bits of the product
- `mulh`: multiply high
 - Gives the upper 64 bits of the product, assuming the operands are signed
- `mulhu`: multiply high unsigned
 - Gives the upper 64 bits of the product, assuming the operands are unsigned
- `mulhsu`: multiply high signed/unsigned
 - Gives the upper 64 bits of the product, assuming one operand is signed and the other unsigned
- Use `mulh` result to check for 64-bit overflow

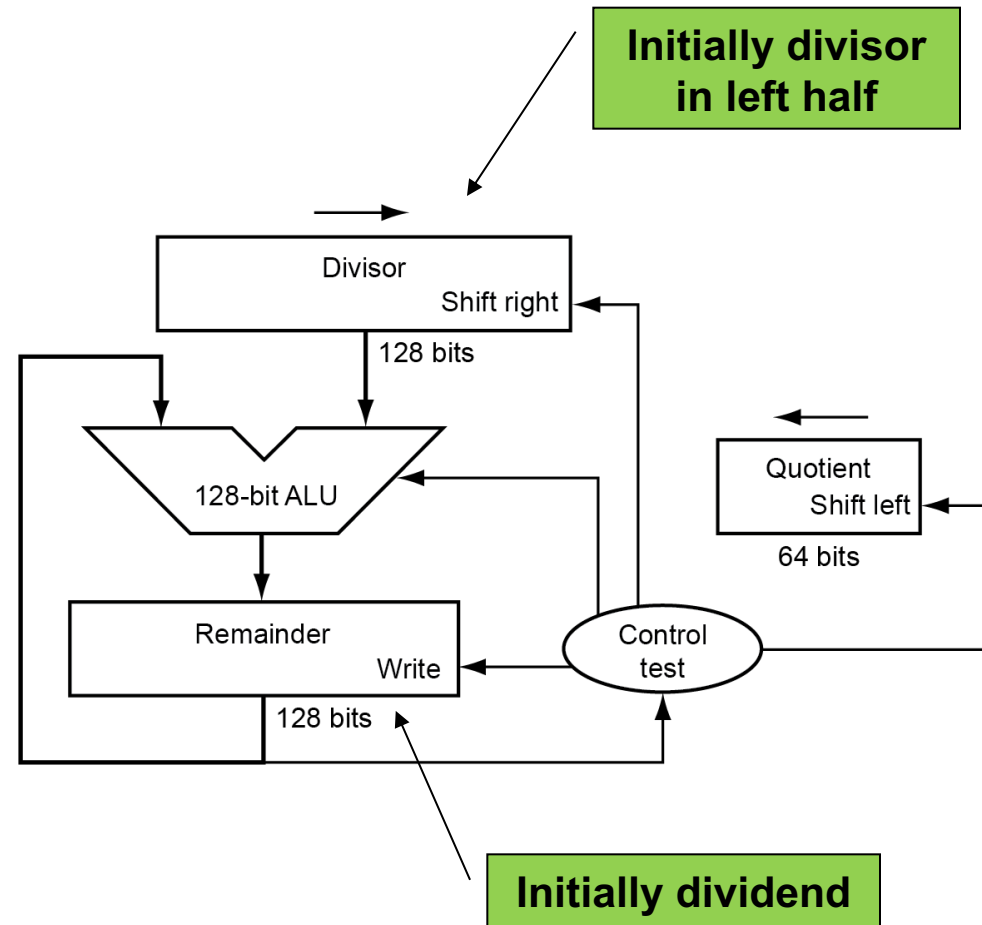
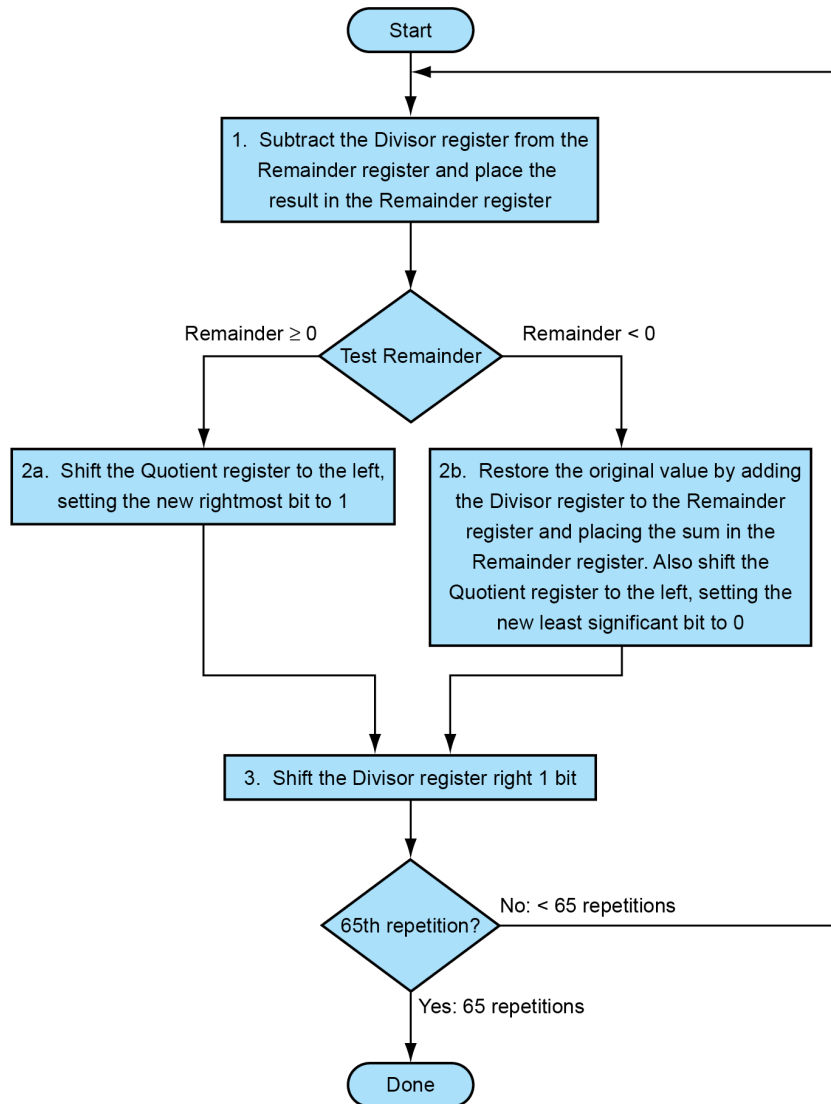
Division



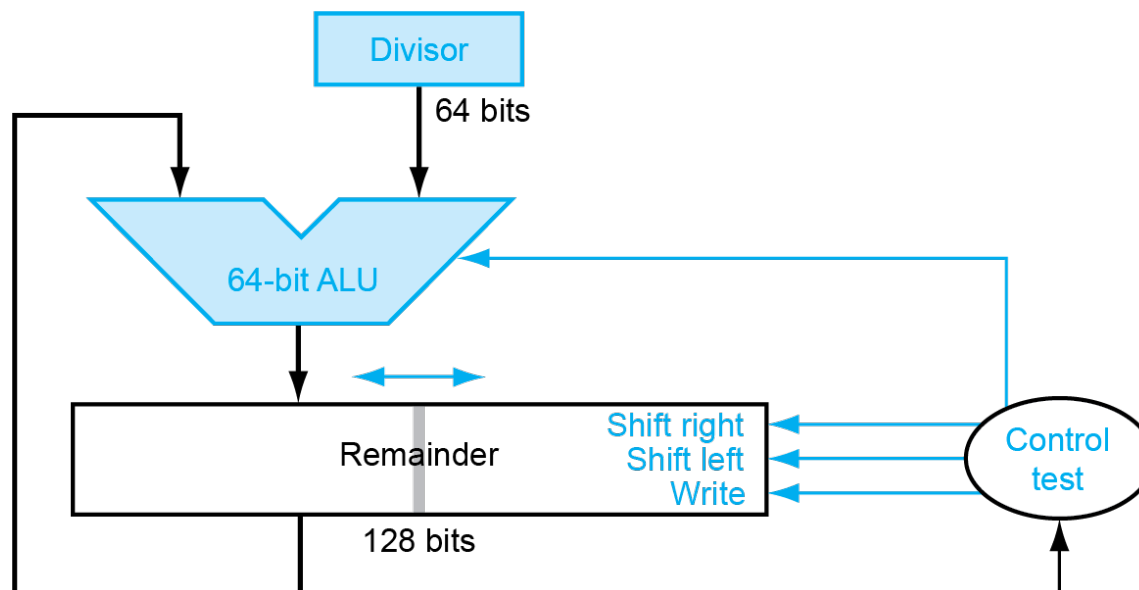
n-bit operands yield *n*-bit quotient and remainder

- **Check for 0 divisor**
- **Long division approach**
 - If divisor \leq dividend bits
 - 1 bit in quotient, subtract
 - Otherwise
 - 0 bit in quotient, bring down next dividend bit
- **Restoring division**
 - Do the subtract, and if remainder goes < 0 , add divisor back
- **Signed division**
 - Divide using absolute values
 - Adjust sign of quotient and remainder as required

Division Hardware



Optimized Divider



- **One cycle per partial-remainder subtraction**
- **Looks a lot like a multiplier!**
 - Same hardware can be used for both

Faster Division

- **Can't use parallel hardware as in multiplier**
 - Subtraction is conditional on sign of remainder
- **Faster dividers (e.g. SRT division) generate multiple quotient bits per step**
 - Still require multiple steps

Right Shift and Division

- Left shift by i places multiplies an integer by 2^i
- Right shift divides by 2^i ?
 - Only for unsigned integers
- For signed integers
 - Arithmetic right shift: replicate the sign bit
 - e.g., $-5 / 4$
 - $11111011_2 \gg 2 = 11111110_2 = -2$
 - Rounds toward $-\infty$
 - c.f. $11111011_2 \ggg 2 = 00111110_2 = +62$

RISC-V Division

■ Four instructions:

- div, rem: signed divide, remainder
- divu, remu: unsigned divide, remainder

■ Overflow and division-by-zero don't produce errors

- Just return defined results
- Faster for the common case of no error

Arithmetic for Computers: An Overview

■ Operations on integers

- Textbook: COD 3.1-3.4
- Addition and subtraction
- Multiplication and division
- Dealing with overflow

■ Floating-point real numbers

- Textbook: P&H 3.5
- Representation and operations

■ Subword parallelism

- Textbook: P&H 3.6

Floating Point

- Representation for non-integral numbers

- Including very small and very large numbers

- Like scientific notation

- -2.34×10^{56}

normalized

- $+0.002 \times 10^{-4}$

- $+987.02 \times 10^9$

not normalized

- In binary

- $\pm 1.xxxxxxx_2 \times 2^{yyyy}$

- Types `float` and `double` in C

Floating Point Standard

- **Defined by IEEE Std 754-1985**
- **Developed in response to divergence of representations**
 - Portability issues for scientific code
- **Now almost universally adopted**
- **Two representations**
 - Single precision (32-bit)
 - Double precision (64-bit)

IEEE Floating-Point Format

single: 8 bits

single: 23 bits

double: 11 bits

double: 52 bits

S	Exponent	Fraction
----------	-----------------	-----------------

$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- **S: sign bit (0 \Rightarrow non-negative, 1 \Rightarrow negative)**
- **Normalize significand: $1.0 \leq |\text{significand}| < 2.0$**
 - Always has a leading pre-binary-point 1 bit, so no need to represent it explicitly (hidden bit)
 - Significand is Fraction with the “1.” restored
- **Exponent: excess representation: actual exponent + Bias**
 - Ensures exponent is unsigned
 - Single: Bias = 127; Double: Bias = 1203

Single-Precision Range

- Exponents 00000000 and 11111111 reserved
- Smallest value
 - Exponent: 00000001
 \Rightarrow actual exponent = $1 - 127 = -126$
 - Fraction: 000...00 \Rightarrow significand = 1.0
 - $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$
- Largest value
 - exponent: 11111110
 \Rightarrow actual exponent = $254 - 127 = +127$
 - Fraction: 111...11 \Rightarrow significand ≈ 2.0
 - $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$

Double-Precision Range

- Exponents 0000...00 and 1111...11 reserved
- Smallest value
 - Exponent: 000000000001
 \Rightarrow actual exponent = $1 - 1023 = -1022$
 - Fraction: 000...00 \Rightarrow significand = 1.0
 - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- Largest value
 - Exponent: 111111111110
 \Rightarrow actual exponent = $2046 - 1023 = +1023$
 - Fraction: 111...11 \Rightarrow significand ≈ 2.0
 - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$

Floating-Point Precision

■ Relative precision

- all fraction bits are significant
- Single: approximately 2^{-23}
 - Equivalent to $23 \times \log_{10} 2 \approx 23 \times 0.3 \approx 6$ decimal digits of precision
- Double: approximately 2^{-52}
 - Equivalent to $52 \times \log_{10} 2 \approx 52 \times 0.3 \approx 16$ decimal digits of precision

Floating-Point Example

■ Represent -0.75

- $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$
- $S = 1$
- Fraction = $1000\dots00_2$
- Exponent = $-1 + \text{Bias}$
 - Single: $-1 + 127 = 126 = 01111110_2$
 - Double: $-1 + 1023 = 1022 = 01111111110_2$

■ Single: **10111111101000...00**

■ Double: **10111111111101000...00**

Floating-Point Example

- What number is represented by the single-precision float

11000000101000...00

- $S = 1$
 - Fraction = $01000...00_2$
 - Exponent = $10000001_2 = 129$
- $x = (-1)^1 \times (1 + 01_2) \times 2^{(129 - 127)}$
 $= (-1) \times 1.25 \times 2^2$
 $= -5.0$

Denormal Numbers

- Exponent = 000...0 \Rightarrow hidden bit is 0

$$x = (-1)^s \times (0 + \text{Fraction}) \times 2^{-\text{Bias}}$$

- Smaller than normal numbers
 - allow for gradual underflow, with diminishing precision
- Denormal with fraction = 000...0

$$x = (-1)^s \times (0 + 0) \times 2^{-\text{Bias}} = \pm 0.0$$

Two representations
of 0.0!



Infinities and NaNs

■ Exponent = 111...1, Fraction = 000...0

- \pm Infinity
- Can be used in subsequent calculations, avoiding need for overflow check

■ Exponent = 111...1, Fraction \neq 000...0

- Not-a-Number (NaN)
- Indicates illegal or undefined result
 - e.g., $0.0 / 0.0$
- Can be used in subsequent calculations

Floating-Point Addition

- **Consider a 4-digit decimal example**

- $9.999 \times 10^1 + 1.610 \times 10^{-1}$

- **1. Align decimal points**

- Shift number with smaller exponent

- $9.999 \times 10^1 + 0.016 \times 10^1$

- **2. Add significands**

- $9.999 \times 10^1 + 0.016 \times 10^1 = 10.015 \times 10^1$

- **3. Normalize result & check for over/underflow**

- 1.0015×10^2

- **4. Round and renormalize if necessary**

- 1.002×10^2

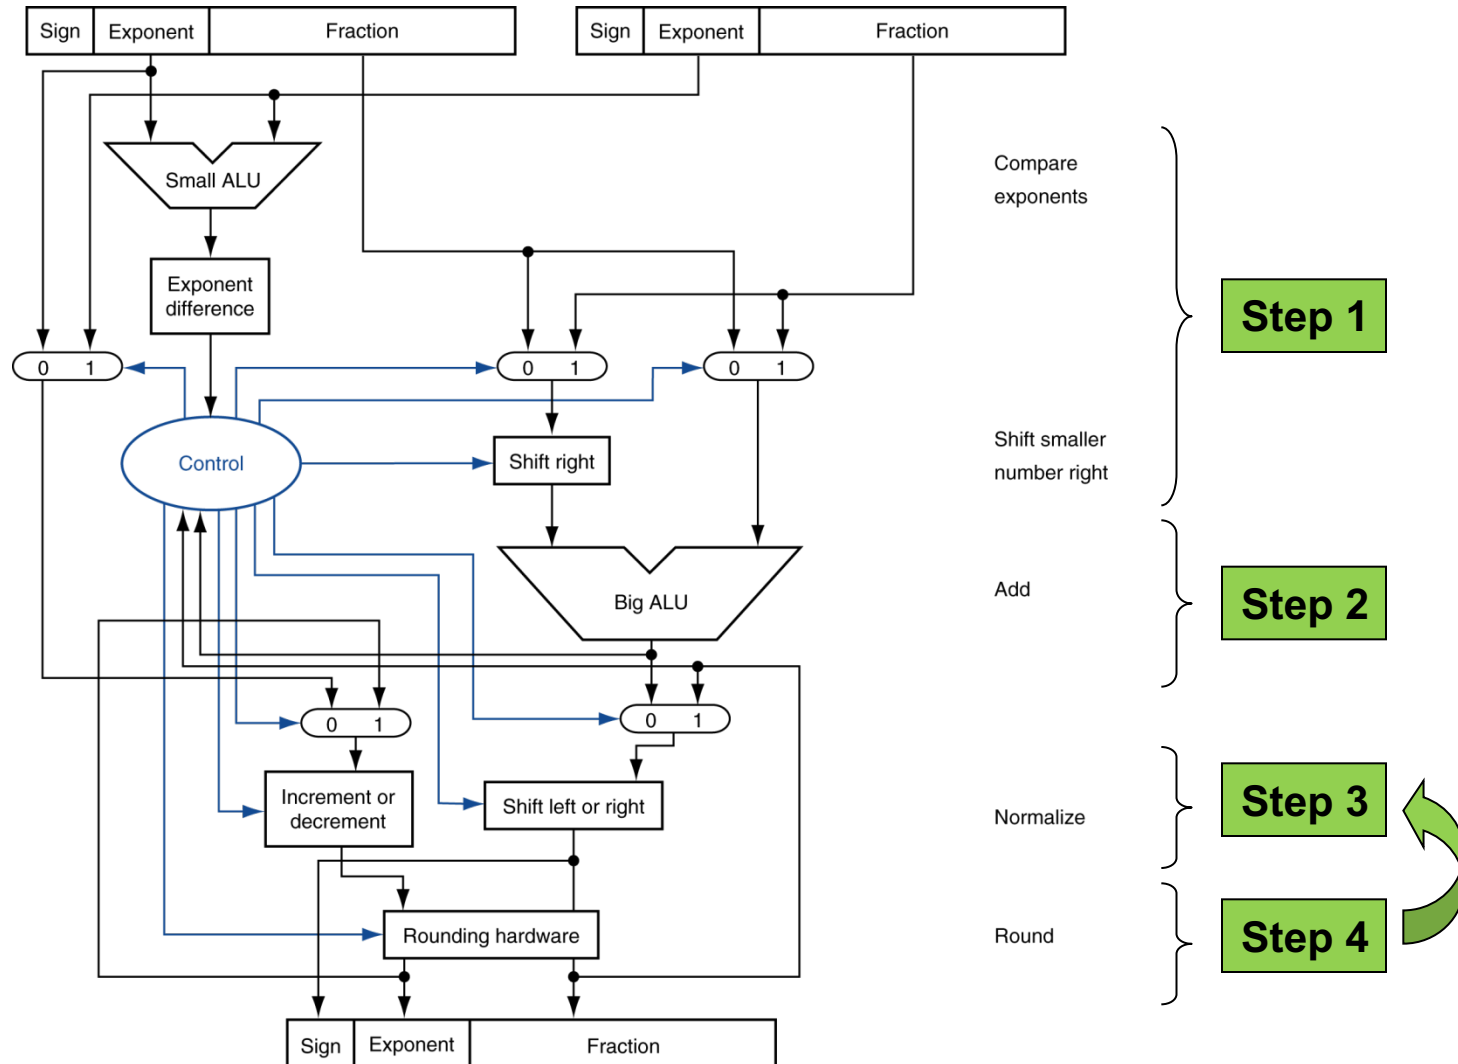
Floating-Point Addition

- **Now consider a 4-digit binary example**
 - $1.000_2 \times 2^{-1} + -1.110_2 \times 2^{-2}$ ($0.5 + -0.4375$)
- **1. Align binary points**
 - Shift number with smaller exponent
 - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1}$
- **2. Add significands**
 - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$
- **3. Normalize result & check for over/underflow**
 - $1.000_2 \times 2^{-4}$, with no over/underflow
- **4. Round and renormalize if necessary**
 - $1.000_2 \times 2^{-4}$ (no change) = 0.0625

FP Adder Hardware

- **Much more complex than integer adder**
- **Doing it in one clock cycle would take too long**
 - Much longer than integer operations
 - Slower clock would penalize all instructions
- **FP adder usually takes several cycles**
 - Can be pipelined

FP Adder Hardware



Floating-Point Multiplication

- **Consider a 4-digit decimal example**
 - $1.110 \times 10^{10} \times 9.200 \times 10^{-5}$
- **1. Add exponents**
 - For biased exponents, subtract bias from sum
 - New exponent = $10 + -5 = 5$
- **2. Multiply significands**
 - $1.110 \times 9.200 = 10.212 \Rightarrow 10.212 \times 10^5$
- **3. Normalize result & check for over/underflow**
 - 1.0212×10^6
- **4. Round and renormalize if necessary**
 - 1.021×10^6
- **5. Determine sign of result from signs of operands**
 - $+1.021 \times 10^6$

Floating-Point Multiplication

■ Now consider a 4-digit binary example

- $1.000_2 \times 2^{-1} \times -1.110_2 \times 2^{-2}$ (0.5×-0.4375)

■ 1. Add exponents

- Unbiased: $-1 + -2 = -3$
- Biased: $(-1 + 127) + (-2 + 127) = -3 + 254 - 127 = -3 + 127$

■ 2. Multiply significands

- $1.000_2 \times 1.110_2 = 1.110_2 \Rightarrow 1.110_2 \times 2^{-3}$

■ 3. Normalize result & check for over/underflow

- $1.110_2 \times 2^{-3}$ (no change) with no over/underflow

■ 4. Round and renormalize if necessary

- $1.110_2 \times 2^{-3}$ (no change)

■ 5. Determine sign: +ve \times -ve \Rightarrow -ve

- $-1.110_2 \times 2^{-3} = -0.21875$

FP Arithmetic Hardware

- **FP multiplier is of similar complexity to FP adder**
 - But uses a multiplier for significands instead of an adder
- **FP arithmetic hardware usually does**
 - Addition, subtraction, multiplication, division, reciprocal, square-root
 - FP \leftrightarrow integer conversion
- **Operations usually takes several cycles**
 - Can be pipelined

FP Instructions in RISC-V

- **Separate FP registers: f0, ..., f31**
 - double-precision
 - single-precision values stored in the lower 32 bits
- **FP instructions operate only on FP registers**
 - Programs generally don't do integer ops on FP data, or vice versa
 - More registers with minimal code-size impact
- **FP load and store instructions**
 - flw, fld
 - fsw, fsd

FP Instructions in RISC-V

■ Single-precision arithmetic

- `fadd.s`, `fsub.s`, `fmul.s`, `fdiv.s`, `fsqrt.s`
 - e.g., `fadds.s f2, f4, f6`

■ Double-precision arithmetic

- `fadd.d`, `fsub.d`, `fmul.d`, `fdiv.d`, `fsqrt.d`
 - e.g., `fadd.d f2, f4, f6`

■ Single- and double-precision comparison

- `feq.s`, `flt.s`, `fle.s`
- `feq.d`, `flt.d`, `fle.d`
- Result is 0 or 1 in integer destination register
 - Use `beq`, `bne` to branch on comparison result

■ Branch on FP condition code true or false

- `B.cond`

FP Example: °F to °C

■ C code:

```
float f2c (float fahr) {  
    return ((5.0/9.0)*(fahr - 32.0));  
}
```

- fahr in f10, result in f10, literals in global memory space

■ Compiled RISC-V code:

f2c:

```
f1w    f0,const5(x3)  // f0 = 5.0f  
f1w    f1,const9(x3)  // f1 = 9.0f  
fdiv.s f0, f0, f1     // f0 = 5.0f / 9.0f  
f1w    f1,const32(x3) // f1 = 32.0f  
fsub.s f10,f10,f1     // f10 = fahr - 32.0  
fmul.s f10,f0,f10     // f10 = (5.0f/9.0f) * (fahr-32.0f)  
jalr   x0,0(x1)       // return
```


FP Example: Array Multiplication

■ $C = C + A \times B$

- All 32×32 matrices, 64-bit double-precision elements

■ C code:

```
void mm(double c[][],
        double a[][], double b[][]) {
    size_t i, j, k;
    for (i = 0; i < 32; i = i + 1)
        for (j = 0; j < 32; j = j + 1)
            for (k = 0; k < 32; k = k + 1)
                c[i][j] = c[i][j]
                    + a[i][k] * b[k][j];
}
```

- Addresses of c, a, b in x10, x11, x12, and
i, j, k in x5, x6, x7

FP Example: Array Multiplication

■ RISC-V code:

mm: . . .

```

        li    x28,32          // x28 = 32 (row size/loop end)
        li    x5,0            // i = 0; initialize 1st for loop
L1:     li    x6,0            // j = 0; initialize 2nd for loop
L2:     li    x7,0            // k = 0; initialize 3rd for loop
        slli  x30,x5,5         // x30 = i * 2**5 (size of row of c)
        add   x30,x30,x6       // x30 = i * size(row) + j
        slli  x30,x30,3        // x30 = byte offset of [i][j]
        add   x30,x10,x30      // x30 = byte address of c[i][j]
        fld   f0,0(x30)       // f0 = c[i][j]
L3:     slli  x29,x7,5         // x29 = k * 2**5 (size of row of b)
        add   x29,x29,x6       // x29 = k * size(row) + j
        slli  x29,x29,3        // x29 = byte offset of [k][j]
        add   x29,x12,x29      // x29 = byte address of b[k][j]
        fld   f1,0(x29)       // f1 = b[k][j]

```

FP Example: Array Multiplication

...

```

slli    x29,x5,5      // x29 = i * 2**5 (size of row of a)
add     x29,x29,x7     // x29 = i * size(row) + k
slli    x29,x29,3     // x29 = byte offset of [i][k]
add     x29,x11,x29    // x29 = byte address of a[i][k]
fld     f2,0(x29)     // f2 = a[i][k]
fmul.d  f1, f2, f1     // f1 = a[i][k] * b[k][j]
fadd.d  f0, f0, f1     // f0 = c[i][j] + a[i][k] * b[k][j]
addi    x7,x7,1       // k = k + 1
bltu    x7,x28,L3     // if (k < 32) go to L3
fsd     f0,0(x30)     // c[i][j] = f0
addi    x6,x6,1       // j = j + 1
bltu    x6,x28,L2     // if (j < 32) go to L2
addi    x5,x5,1       // i = i + 1
bltu    x5,x28,L1     // if (i < 32) go to L1

```

Accurate Arithmetic

- **IEEE Std 754 specifies additional rounding control**
 - Extra bits of precision (guard, round, sticky)
 - Choice of rounding modes
 - Allows programmer to fine-tune numerical behavior of a computation
- **Not all FP units implement all options**
 - Most programming languages and FP libraries just use defaults
- **Trade-off between hardware complexity, performance, and market requirements**

Associativity

- Parallel programs may interleave operations in unexpected orders
 - Assumptions of associativity may fail

		$(x+y)+z$	$x+(y+z)$
x	-1.50E+38		-1.50E+38
y	1.50E+38	0.00E+00	
z	1.0	1.0	1.50E+38
		1.00E+00	0.00E+00

- Need to validate parallel programs under varying degrees of parallelism

Who Cares About FP Accuracy?

- **Important for scientific code**
 - But for everyday consumer use?
 - “My bank balance is out by 0.0002¢!” ☹️
- **The Intel Pentium FDIV bug**
 - The market expects accuracy
 - See Colwell, The Pentium Chronicles

Arithmetic for Computers: An Overview

■ Operations on integers

- Textbook: COD 3.1-3.4
- Addition and subtraction
- Multiplication and division
- Dealing with overflow

■ Floating-point real numbers

- Textbook: P&H 3.5
- Representation and operations

■ Subword parallelism

- Textbook: P&H 3.6


Subword Parallellism

- **Graphics and audio applications can take advantage of performing simultaneous operations on short vectors**
 - Example: 128-bit adder:
 - Sixteen 8-bit adds
 - Eight 16-bit adds
 - Four 32-bit adds
- **Also called data-level parallelism, vector parallelism, or Single Instruction, Multiple Data (SIMD)**

SIMD Extensions

■ Example: Intel MMX (1997)

- 57 additional 80x86 instructions (1st since 386)
 - similar to Intel 860, Mot. 88110, HP PA-71000LC, UltraSPARC
- 3 data types: 8 8-bit, 4 16-bit, 2 32-bit in 64bits
 - reuse 8 FP registers (FP and MMX cannot mix)
- short vector: load, add, store 8 8-bit operands



- Claim: overall speedup 1.5 to 2X for 2D/3D graphics, audio, video, speech, comm., ...
 - use in drivers or added to library routines; no compiler

SIMD Extensions

■ Example: Intel MMX instructions

- Move 32b, 64b
- Add, Subtract in parallel: 8 8b, 4 16b, 2 32b
 - opt. signed/unsigned saturate (set to max) if overflow
- Shifts (sll,srl, sra), And, And Not, Or, Xor in parallel: 8 8b, 4 16b, 2 32b
- Multiply, Multiply-Add in parallel: 4 16b
- Compare = , > in parallel: 8 8b, 4 16b, 2 32b
 - sets field to 0s (false) or 1s (true); removes branches
- Pack/Unpack
 - Convert 32b \leftrightarrow 16b, 16b \leftrightarrow 8b
 - Pack saturates (set to max) if number is too large

SIMD Extensions

■ Streaming SIMD Extension 2 (SSE2) (2001)

- Adds 4×128 -bit registers
 - Extended to 8 registers in AMD64/EM64T
- Can be used for multiple FP operands
 - 2×64 -bit double precision
 - 4×32 -bit double precision
 - Instructions operate on them simultaneously
 - Single-Instruction Multiple-Data

■ Up-to-date version: AVX-512 (2015)

- Advanced Vector Extensions
- 512-bit SIMD registers
- Targets machine learning, high-performance computing, etc.

Concluding Remarks

- **Bits have no inherent meaning**
 - Interpretation depends on the instructions applied
- **Computer representations of numbers**
 - Finite range and precision
 - Need to account for this in programs