


Jae W. Lee (jaewlee@snu.ac.kr), SNU Computer Science and Engineering

Slide credits: [CS:APP3e] slides from CMU; [COD:RV2e] slides from Elsevier Inc.

The Processor (2)

Lecture 5

October 4th, 2023



CompArch Today #2: [Apple introduces M2 Ultra] (06/05/2023)

<https://www.apple.com/newsroom/2023/06/apple-introduces-m2-ultra/>

M2 Ultra is built using a second-generation 5-nanometer process and uses Apple's UltraFusion technology to connect the die of two M2 Max chips. UltraFusion uses a silicon interposer that connects the dies with more than 10,000 signals, providing over 2.5TB/s of low-latency interprocessor bandwidth. UltraFusion's architecture enables M2 Ultra to appear as a single chip to software. It has 134 billion transistors, 20 billion more than the M1 Ultra. It also has a unified memory architecture that supports up to 192GB of memory capacity and 800GB/s of memory bandwidth. The M2 Ultra features a more powerful 24-core CPU (16 performance + 8 efficiency cores) that's 20 percent faster than the M1 Ultra, a larger GPU that's up to 30 percent faster, and a Neural Engine (with 31.6 trillion operations per second (TOPS)) that's up to 40 percent faster.


Outline

Textbook: P&H 4.5-4.6

- **An Overview of Pipelining**
- **Pipelined Datapath and Control**


You Already Know Pipelining: Laundry Example

■ Sequential Processing: Wash-Dry-Fold-Store

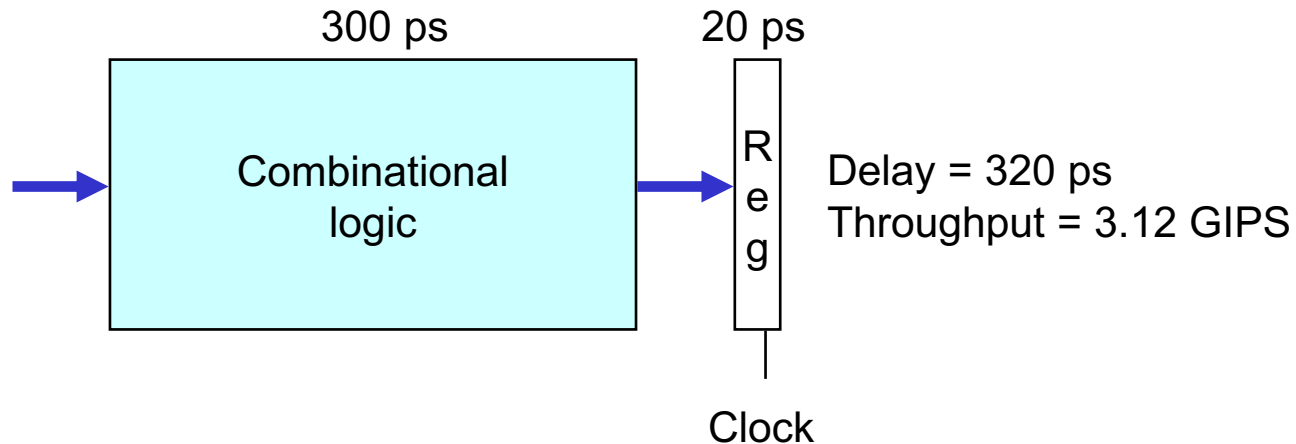


You Already Know Pipelining: Laundry Example

■ Pipelined Processing



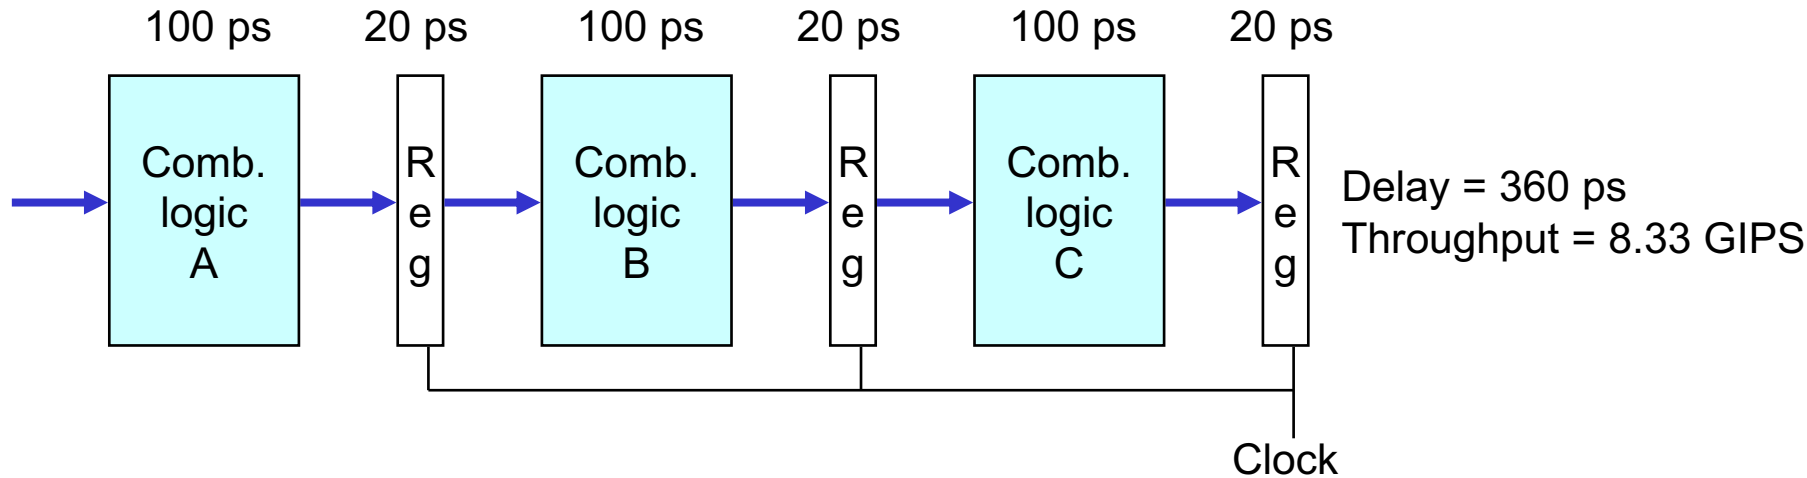
Pipelining for Computation



■ System

- Computation requires total of 300 picoseconds
- Additional 20 picoseconds to save result in register
- Must have clock cycle of at least 320 ps

Pipelining for Computation




■ 3-Way Pipelined Version

- Divide combinational logic into 3 blocks of 100 ps each
- Can begin new operation as soon as previous one passes through stage A.
 - Begin new operation every 120 ps
- Overall latency increases
 - 360 ps from start to finish


Pipelining for Computation: Pipeline Diagrams

■ Unpipelined



- Cannot start new operation until previous one completes

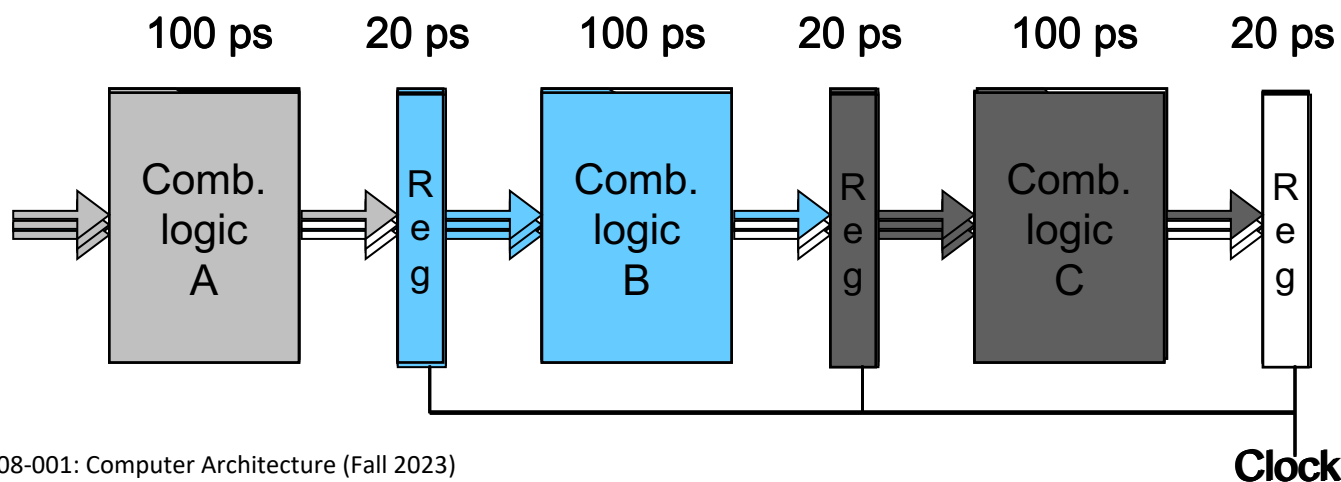
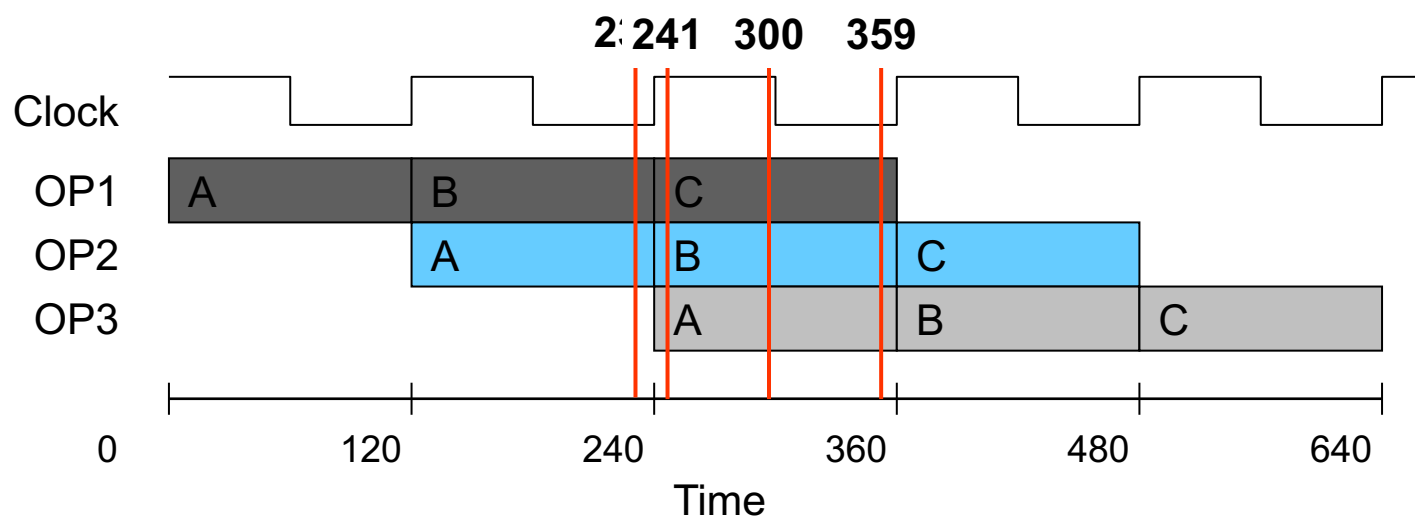
■ 3-Way Pipelined



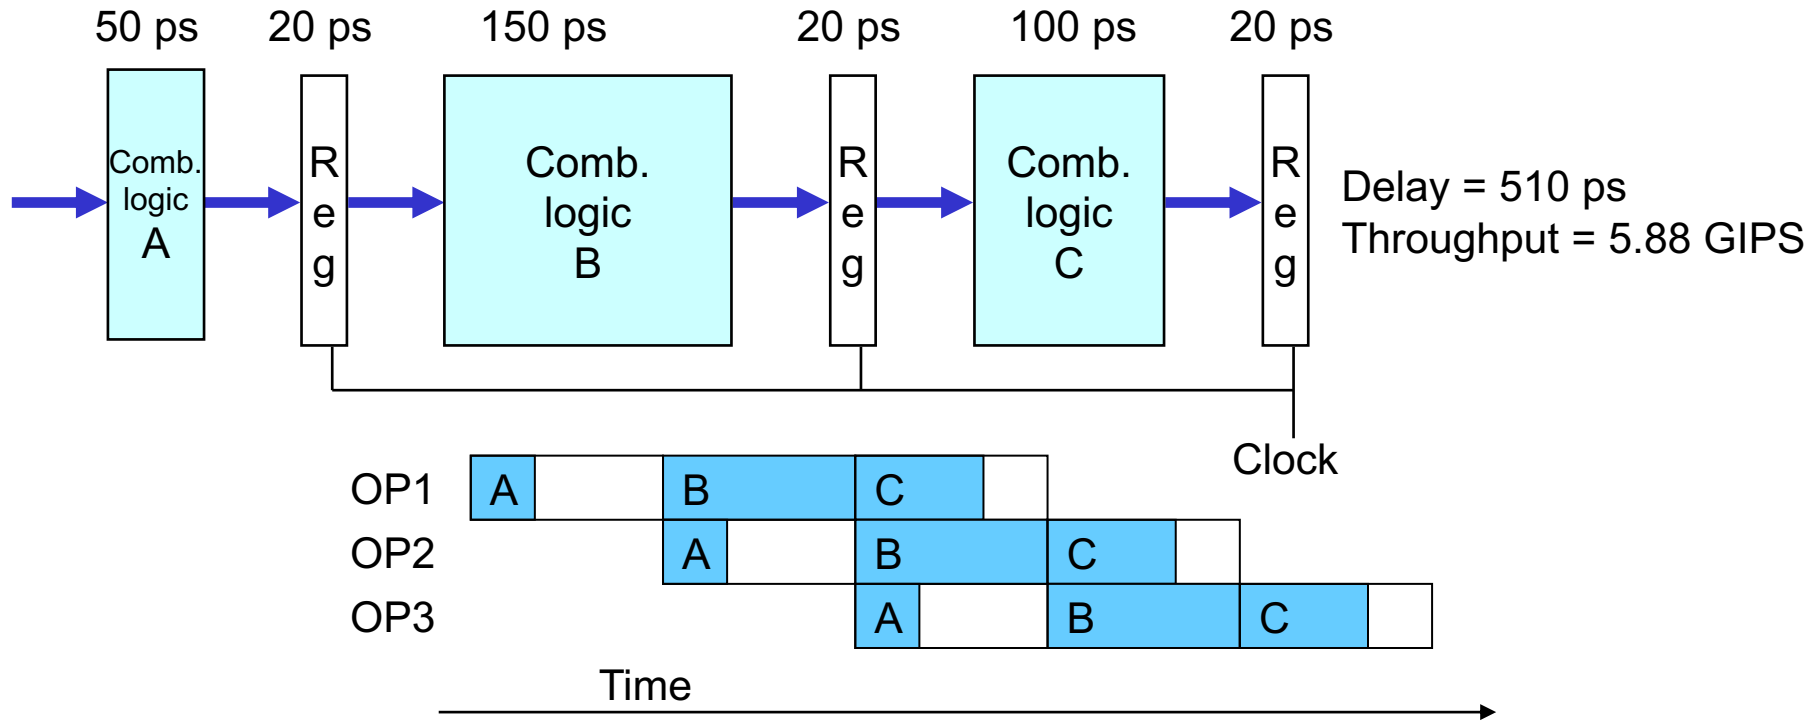
- Up to 3 operations in process simultaneously

Pipelining for Computation

■ 3-Way Pipelined Version: Operation

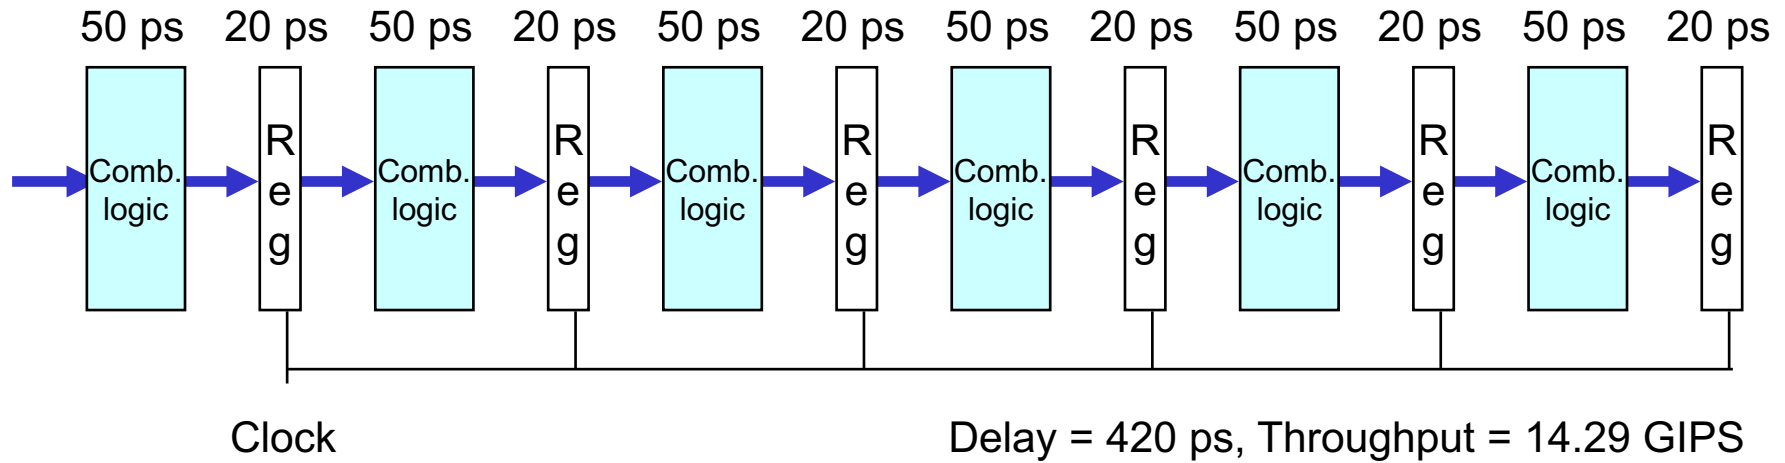


Limitations of Pipelining: Nonuniform Delays



- Throughput limited by slowest stage
- Other stages sit idle for much of the time
- Challenging to partition system into balanced stages

Limitations of Pipelining: Register Overhead



- As try to deepen pipeline, overhead of loading registers becomes more significant
- Percentage of clock cycle spent loading register:
 - 1-stage pipeline: 6.25%
 - 3-stage pipeline: 16.67%
 - 6-stage pipeline: 28.57%
- High speeds of modern processor designs obtained through very deep pipelining

RISC-V Pipeline

- **Five stages, one step per stage**
 - IF: Instruction fetch from memory
 - ID: Instruction decode & register read
 - EX: Execute operation or calculate address
 - MEM: Access memory operand
 - WB: Write result back to register

Pipelined Instruction Execution

■ Sequential Execution



■ Pipelined Execution



add x21, x12, x0



sub x22, x21, x20



and x16, x22, 0xFFF

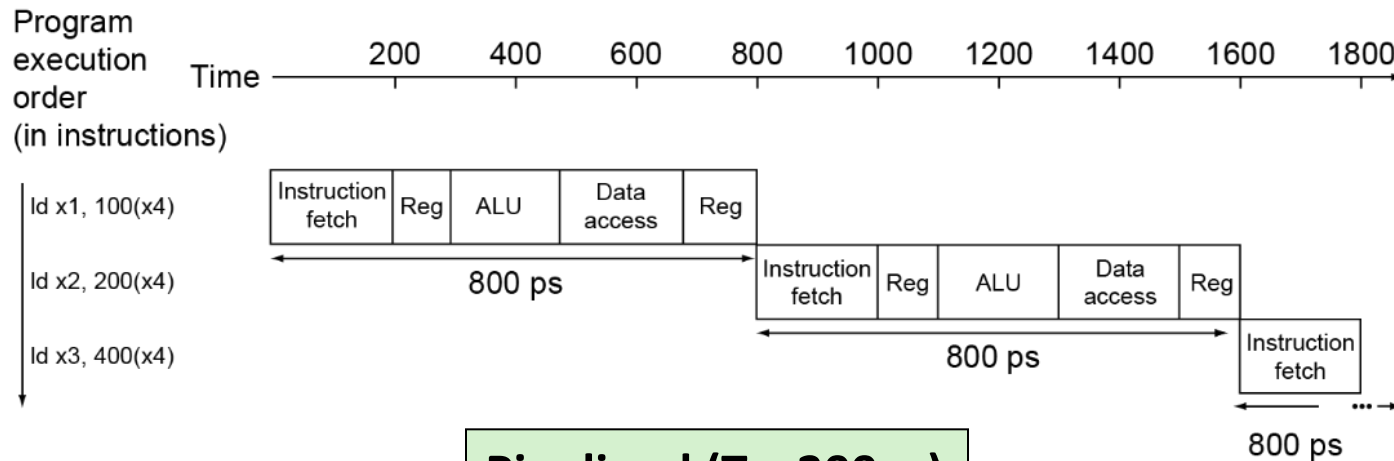
Pipeline Performance

- Assume time for stages is
 - 100ps for register read or write
 - 200ps for other stages
- Compare pipelined datapath with single-cycle datapath

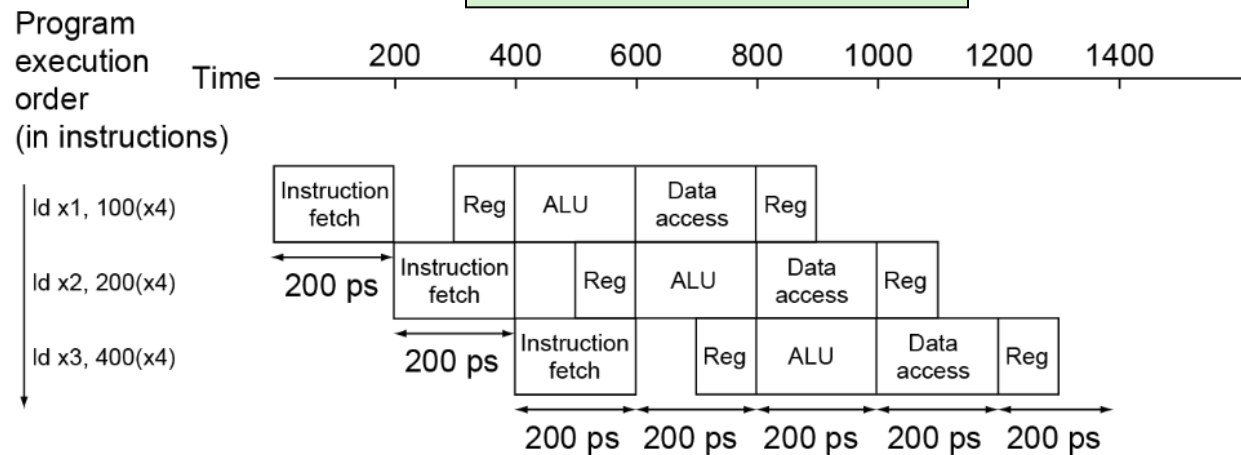
Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
ld	200ps	100 ps	200ps	200ps	100 ps	800ps
sd	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

Pipeline Performance

Single-cycle ($T_c = 800\text{ps}$)



Pipelined ($T_c = 200\text{ps}$)



Pipeline Speedup

- **If all stages are balanced**

- i.e., all take the same time
- Time between instructions_{pipelined}
= $\frac{\text{Time between instructions}_{\text{nonpipelined}}}{\text{Number of stages}}$

- **If not balanced, speedup is less**

- **Speedup due to increased throughput**

- Latency (time for each instruction) does not decrease

Pipelining and ISA Design

■ RISC-V ISA designed for pipelining

- All instructions are 32-bits
 - Easier to fetch and decode in one cycle
 - *cf.* x86: 1- to 17-byte instructions
- Few and regular instruction formats
 - Can decode and read registers in one step
- Load/store addressing
 - Can calculate address in 3rd stage, access memory in 4th stage

Hazards: Major Hurdles of Pipelining

- Situations that prevent starting the next instruction in the next cycle
- Structure hazards
 - A required resource is busy
- Data hazard
 - Need to wait for previous instruction to complete its data read/write
- Control hazard
 - Deciding on control action depends on previous instruction

Structure Hazards

- **Conflict for use of a resource**
- **Example: in (hypothetical) RISC-V pipeline with a single memory**
 - Load/store requires data access
 - Instruction fetch would have to stall for that cycle
 - Would cause a pipeline “bubble”

Instruction number	Clock number								
	1	2	3	4	5	6	7	8	9
Load Instruction	F	D	E	M	W				
Instruction $i + 1$		F	D	E	M	W			
Instruction $i + 2$			F	D	E	M	W		
Instruction $i + 3$				F	D	E	M	W	
Instruction $i + 4$					F	D	E	M	W

Structure Hazards

■ Solutions to Structural Hazard: Resource Duplication

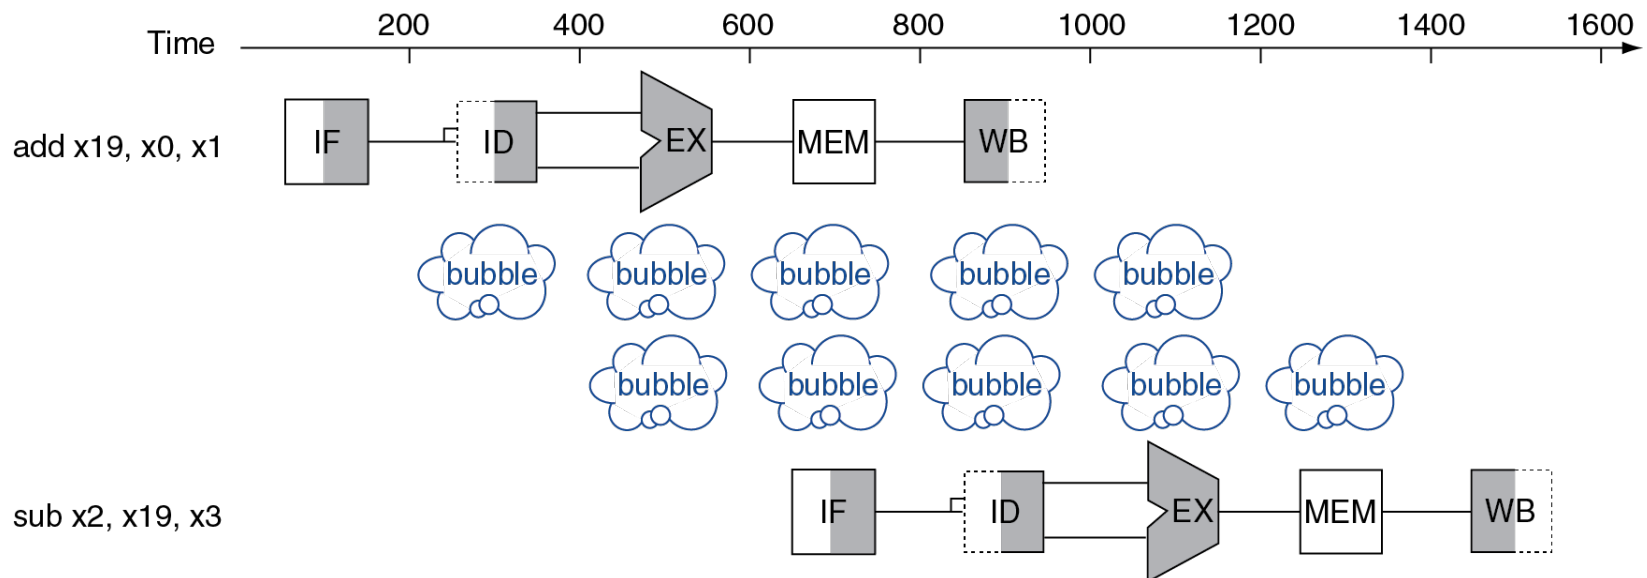
■ example

- Separate I and D caches for memory access conflict
- Multi-port register file for register file access conflict

Data Hazards

- An instruction depends on completion of data access by a previous instruction

```
add    x19, x0, x1  
sub     x2, x19, x3
```



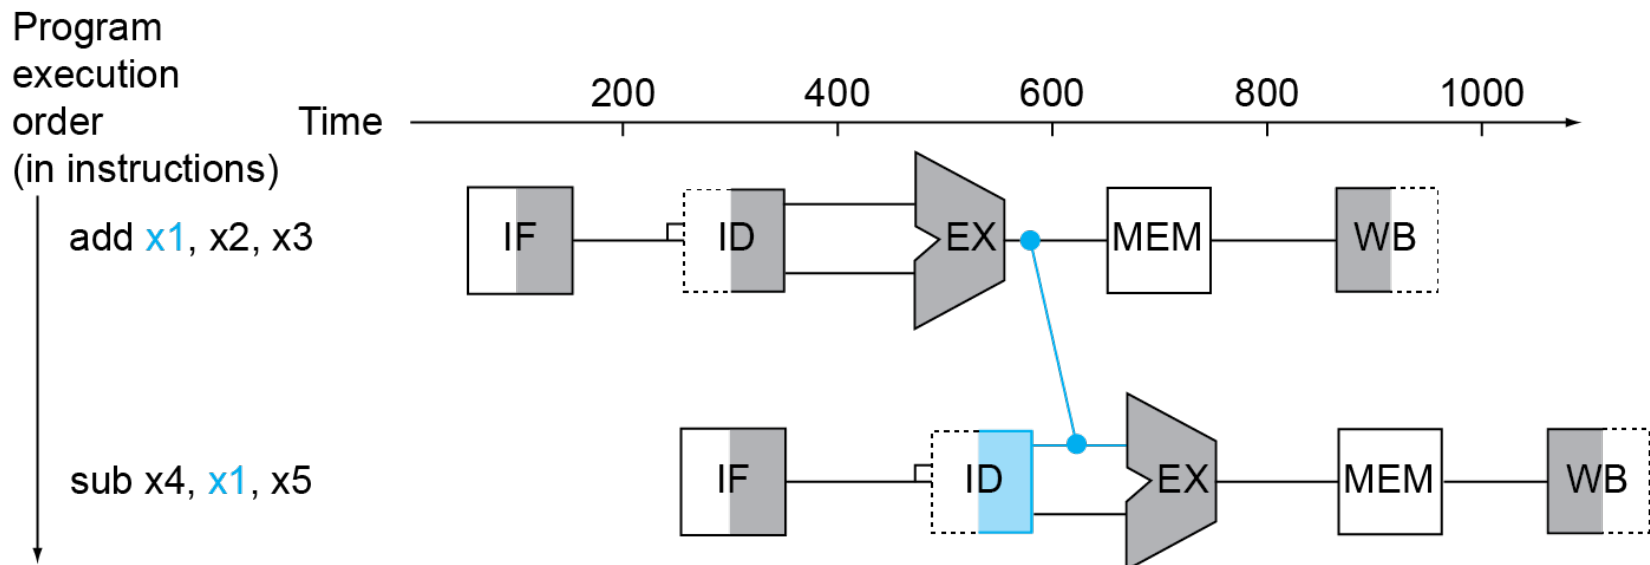
Data Hazards

■ Solutions to Data Hazard

1. Freezing the pipeline
2. (Internal) Forwarding
3. Compiler scheduling

Forwarding (aka Bypassing)

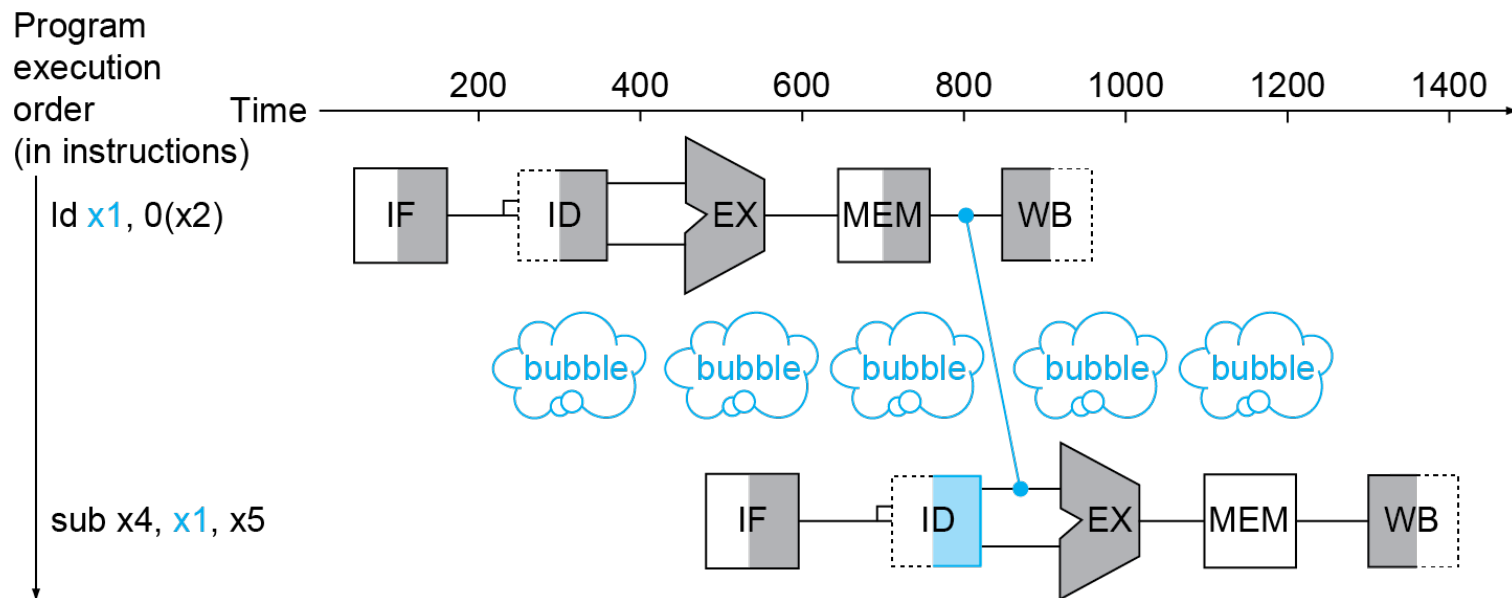
- Use result when it is computed
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath



Load-Use Data Hazard

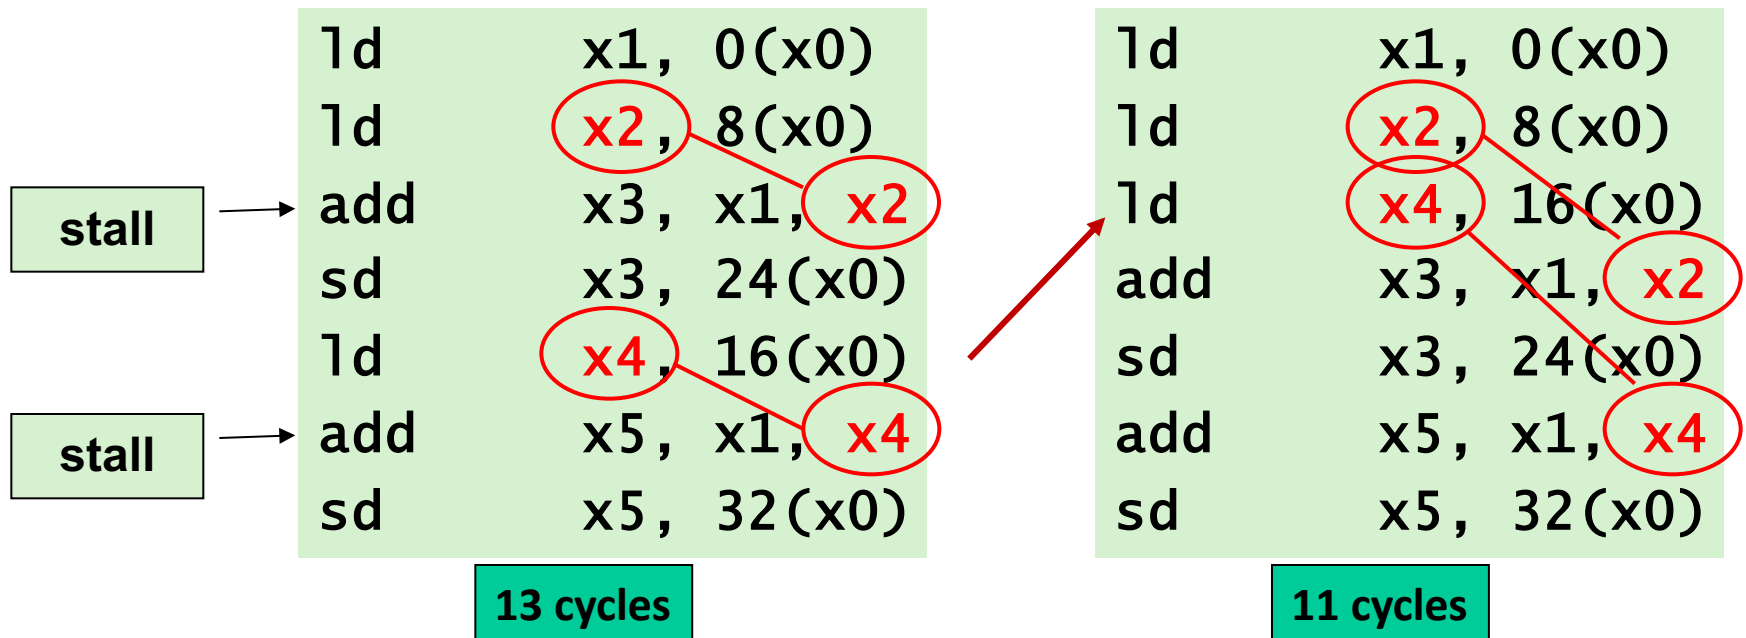
■ Can't always avoid stalls by forwarding

- If value not computed when needed
- Can't forward backward in time!



Compiler Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction
- C code for `a = b + e; c = b + f;`



Control Hazards

■ Branch determines flow of control

- Fetching next instruction depends on branch outcome
- Pipeline can't always fetch correct instruction
 - Still working on ID stage of branch

(Example)

Branch Instruction	F	D	E	M	W					
Branch successor		F	<i>stall</i>	<i>stall</i>	F	D	E	M	W	
Branch successor + 1						F	D	E	M	W
Branch successor + 2							F	D	E	M
Branch successor + 3								F	D	E
Branch successor + 4									F	D
Branch successor + 5										F

For 5-stage pipeline, 3 cycle penalty
15% branch frequency. CPI = 1.45

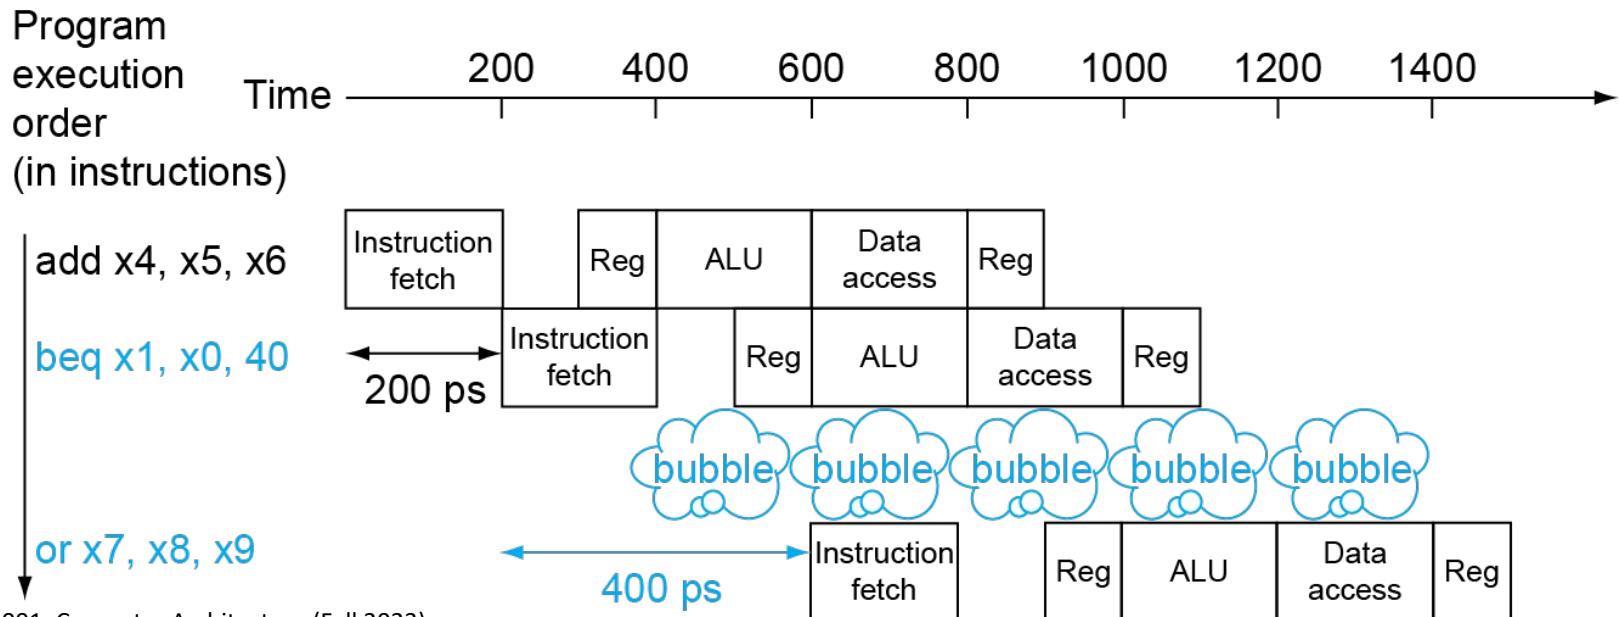
Control Hazards

■ Solution #1: Stall on branch

- Wait until branch outcome determined before fetching next instruction

■ In RISC-V pipeline

- Need to compare registers and compute target early in the pipeline
- Add hardware to do it in ID stage



Control Hazards

■ Solution #2: Branch prediction

- Longer pipelines can't readily determine branch outcome early
 - Stall penalty becomes unacceptable
- Predict outcome of branch
 - Only stall if prediction is wrong

■ In RISC-V pipeline

- Can predict branches not taken
- Fetch instruction after branch, with no delay

More-Realistic Branch Prediction

■ Static branch prediction

- Based on typical branch behavior
- Example: loop and if-statement branches
 - Predict backward branches taken
 - Predict forward branches not taken

■ Dynamic branch prediction

- Hardware measures actual branch behavior
 - e.g., record recent history of each branch
- Assume future behavior will continue the trend
 - When wrong, stall while re-fetching, and update history

Pipeline Summary

- **Pipelining improves performance by increasing instruction throughput**
 - Executes multiple instructions in parallel
 - Each instruction has the same latency
- **Subject to hazards**
 - Structure, data, control
- **Instruction set design affects complexity of pipeline implementation**

Outline

Textbook: P&H 4.5-4.6

- **An Overview of Pipelining**
- **Pipelined Datapath and Control**

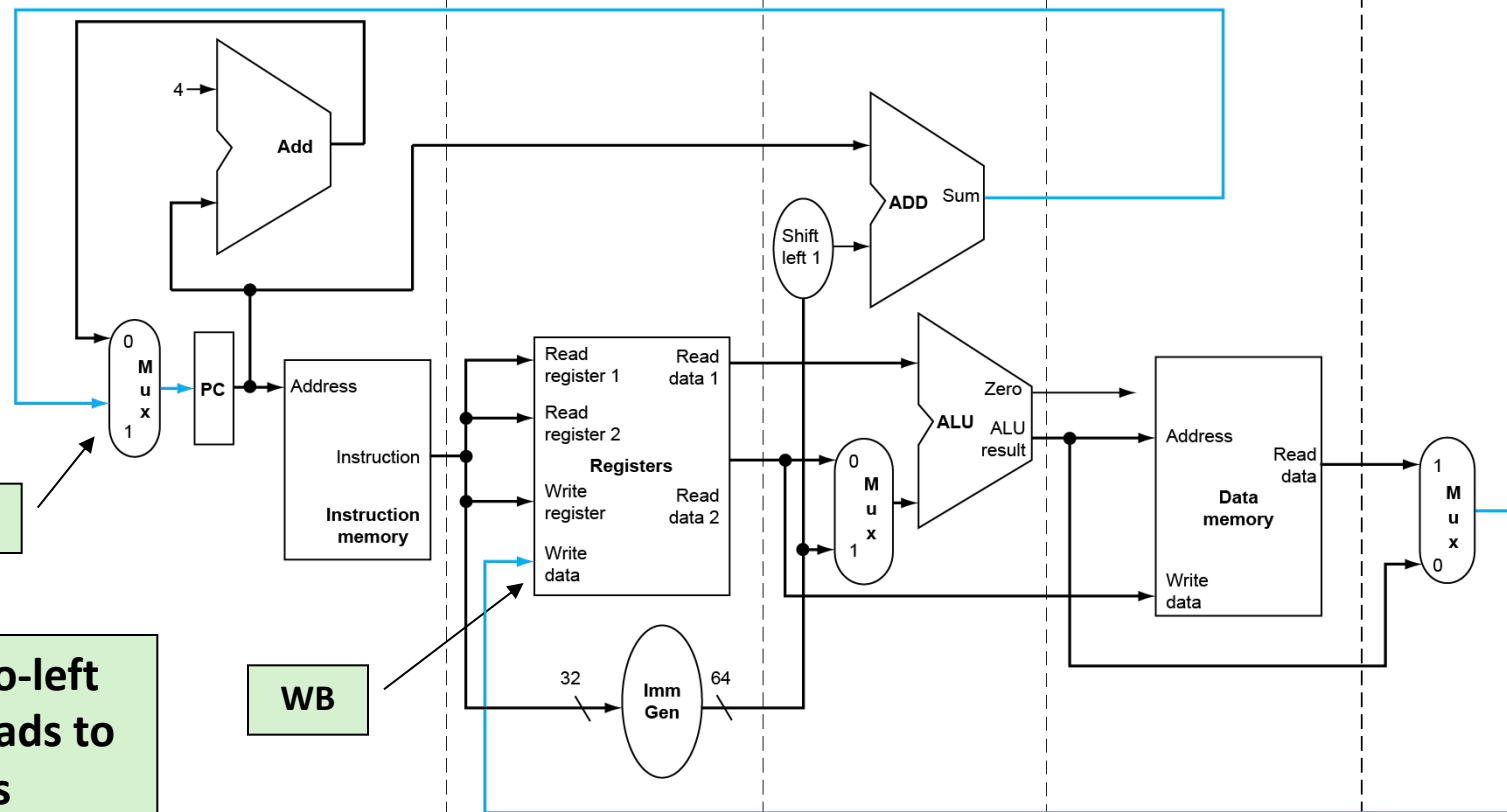
RISC-V Pipelined Datapath

IF: Instruction fetch

ID: Instruction decode/
register file readEX: Execute/
address calculation

MEM: Memory access

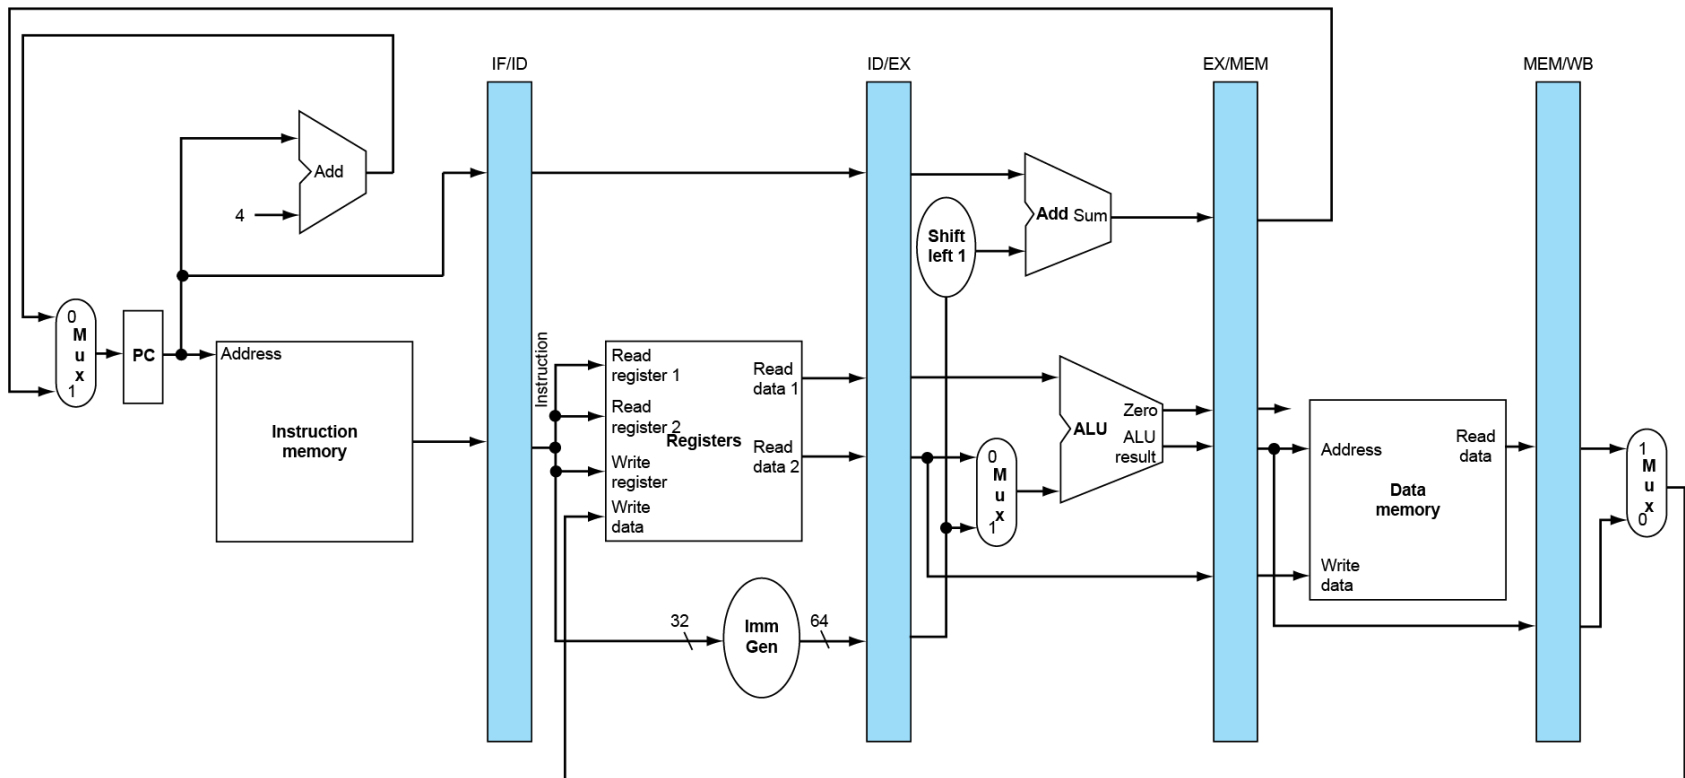
WB: Write back



Right-to-left
flow leads to
hazards

Pipeline registers

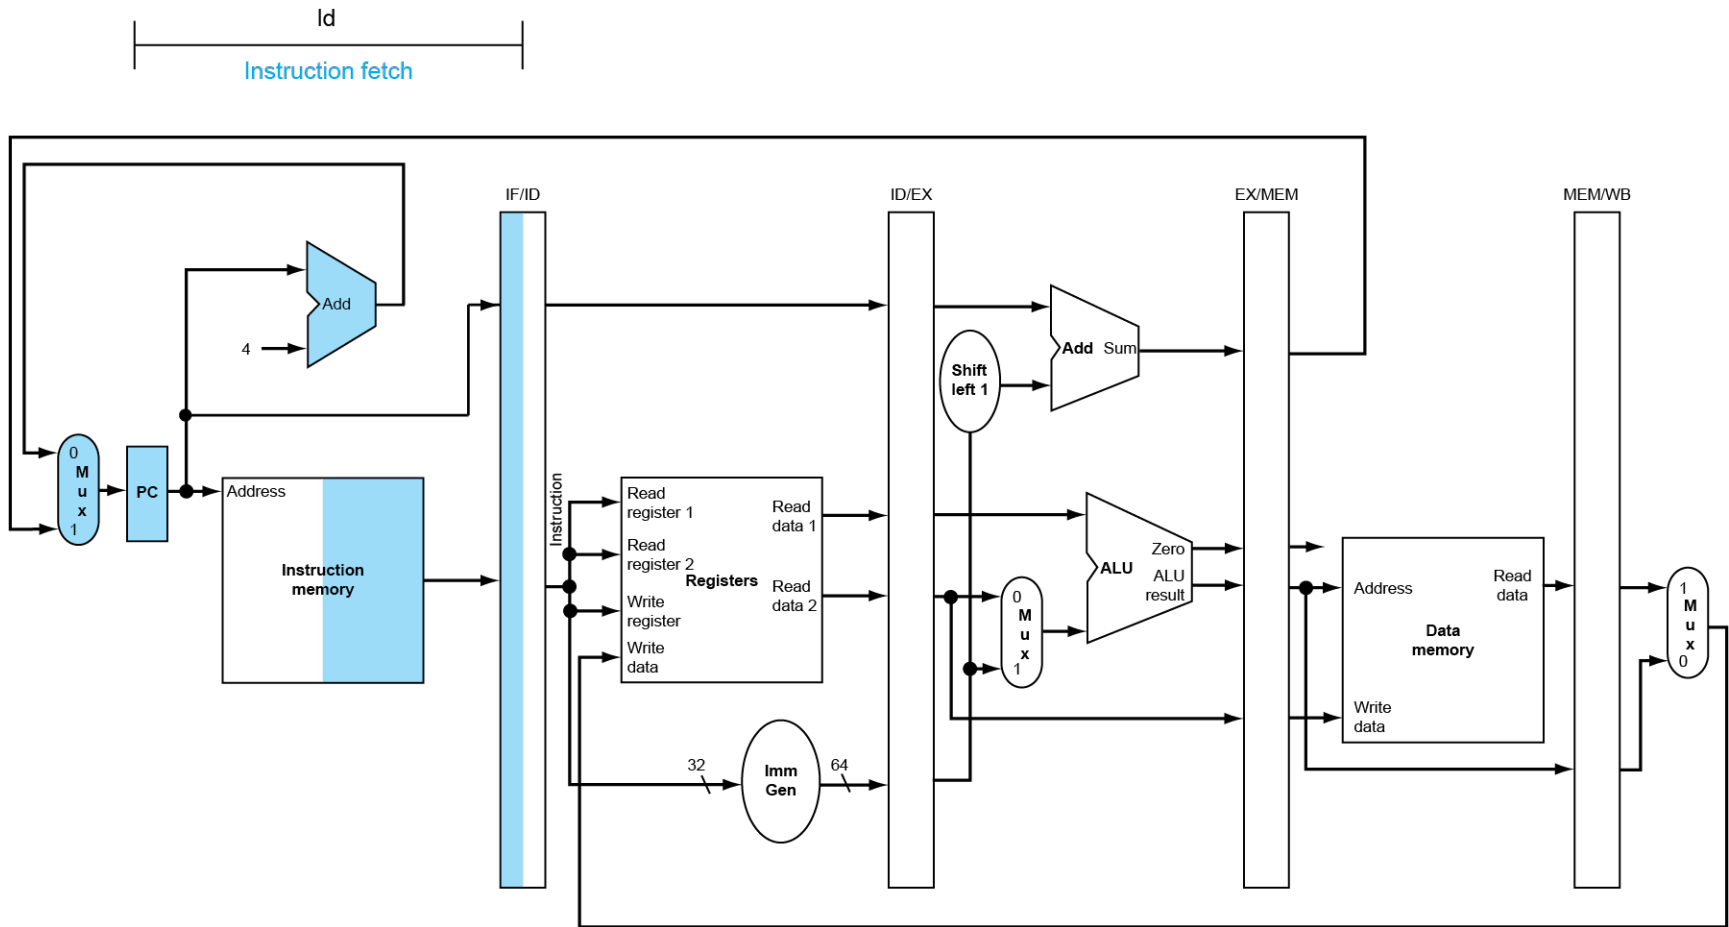
- Need registers between stages
 - To hold information produced in previous cycle



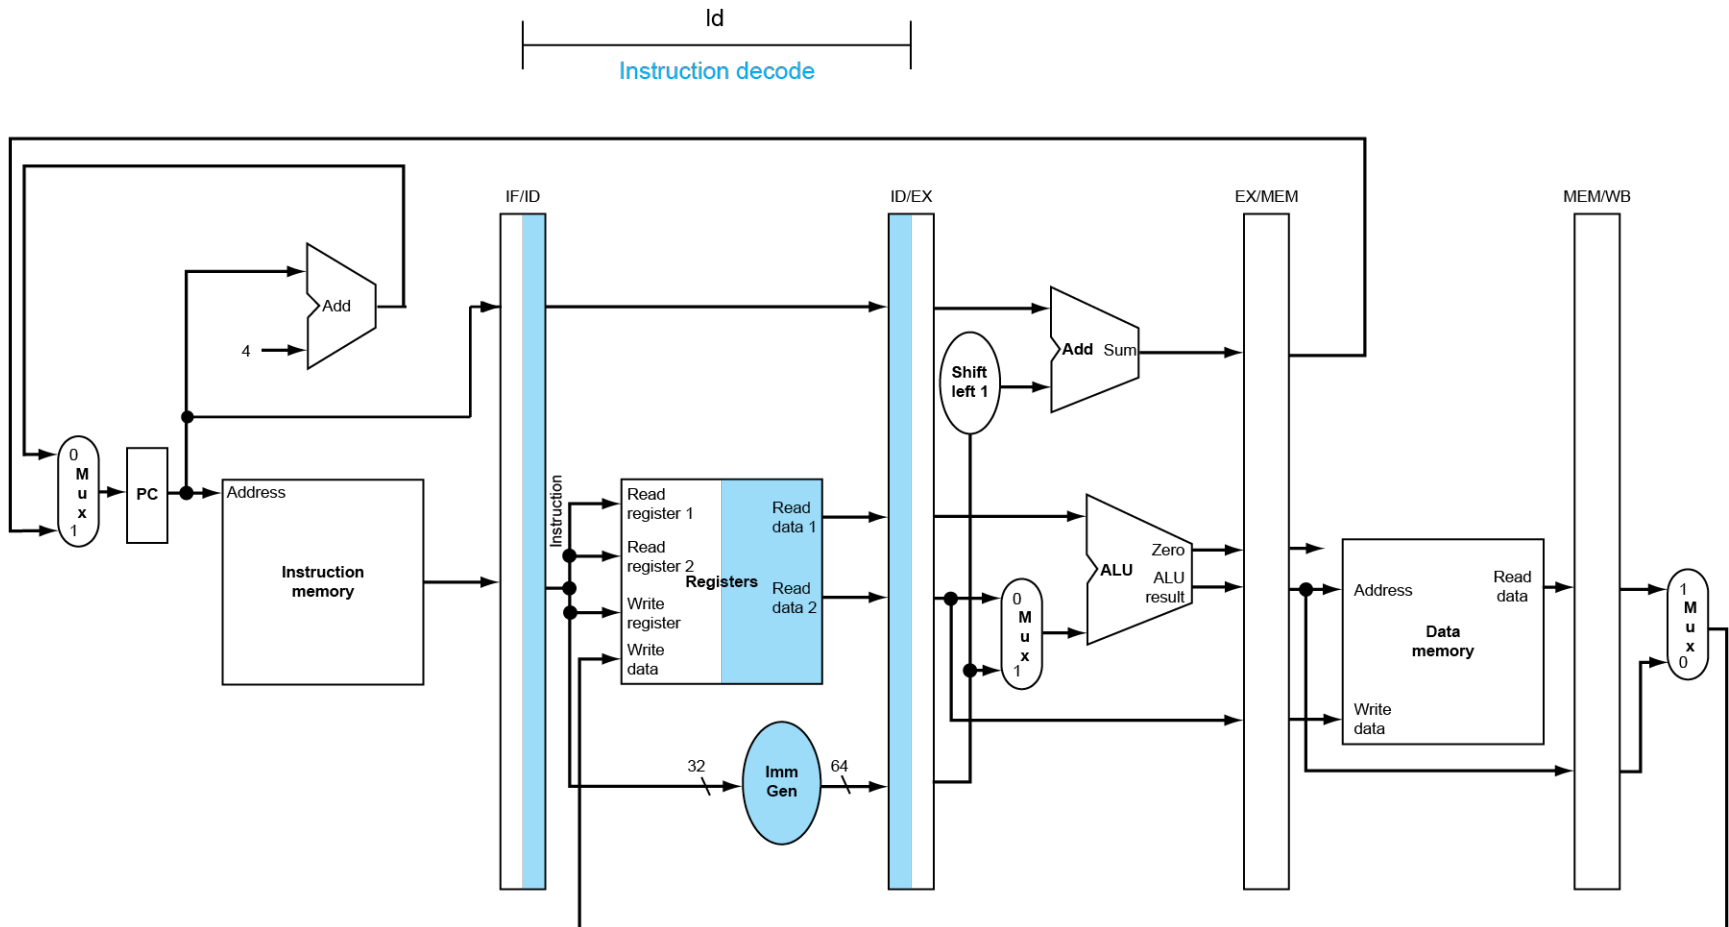
Pipeline Operation

- **Cycle-by-cycle flow of instructions through the pipelined datapath**
 - “Single-clock-cycle” pipeline diagram
 - Shows pipeline usage in a single cycle
 - Highlight resources used
 - *cf.* “multi-clock-cycle” diagram
 - Graph of operation over time
- **We’ll look at “single-clock-cycle” diagrams for load & store**

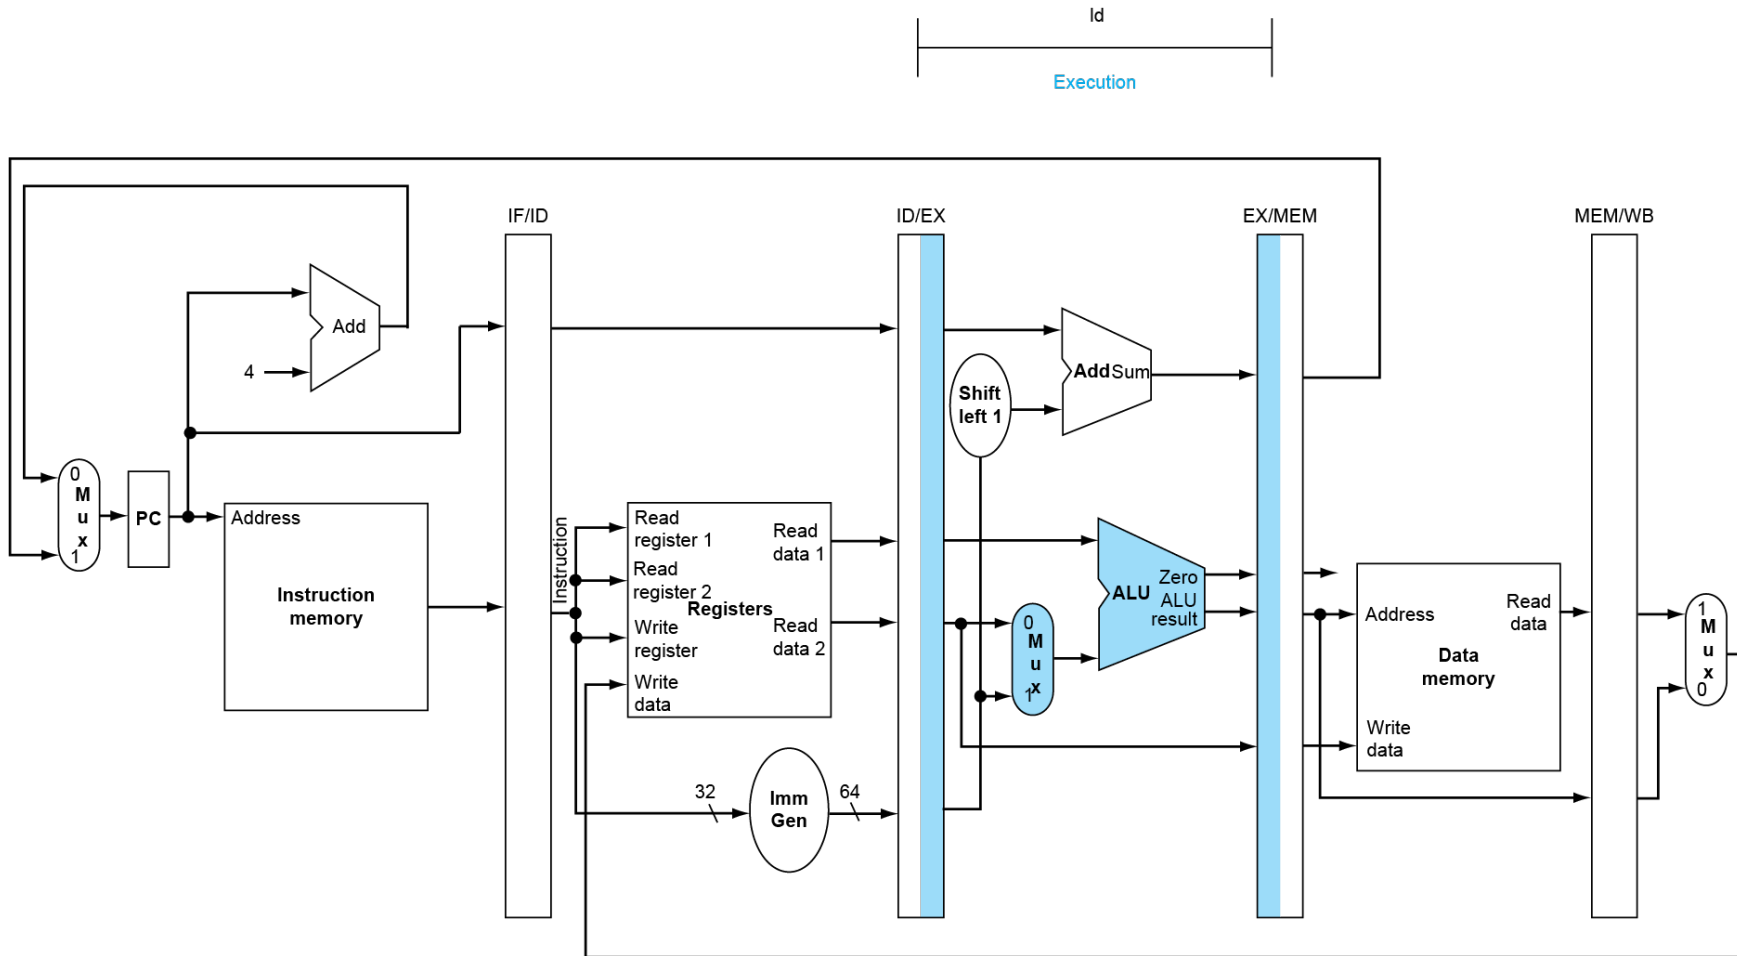
IF for Load, Store, ...



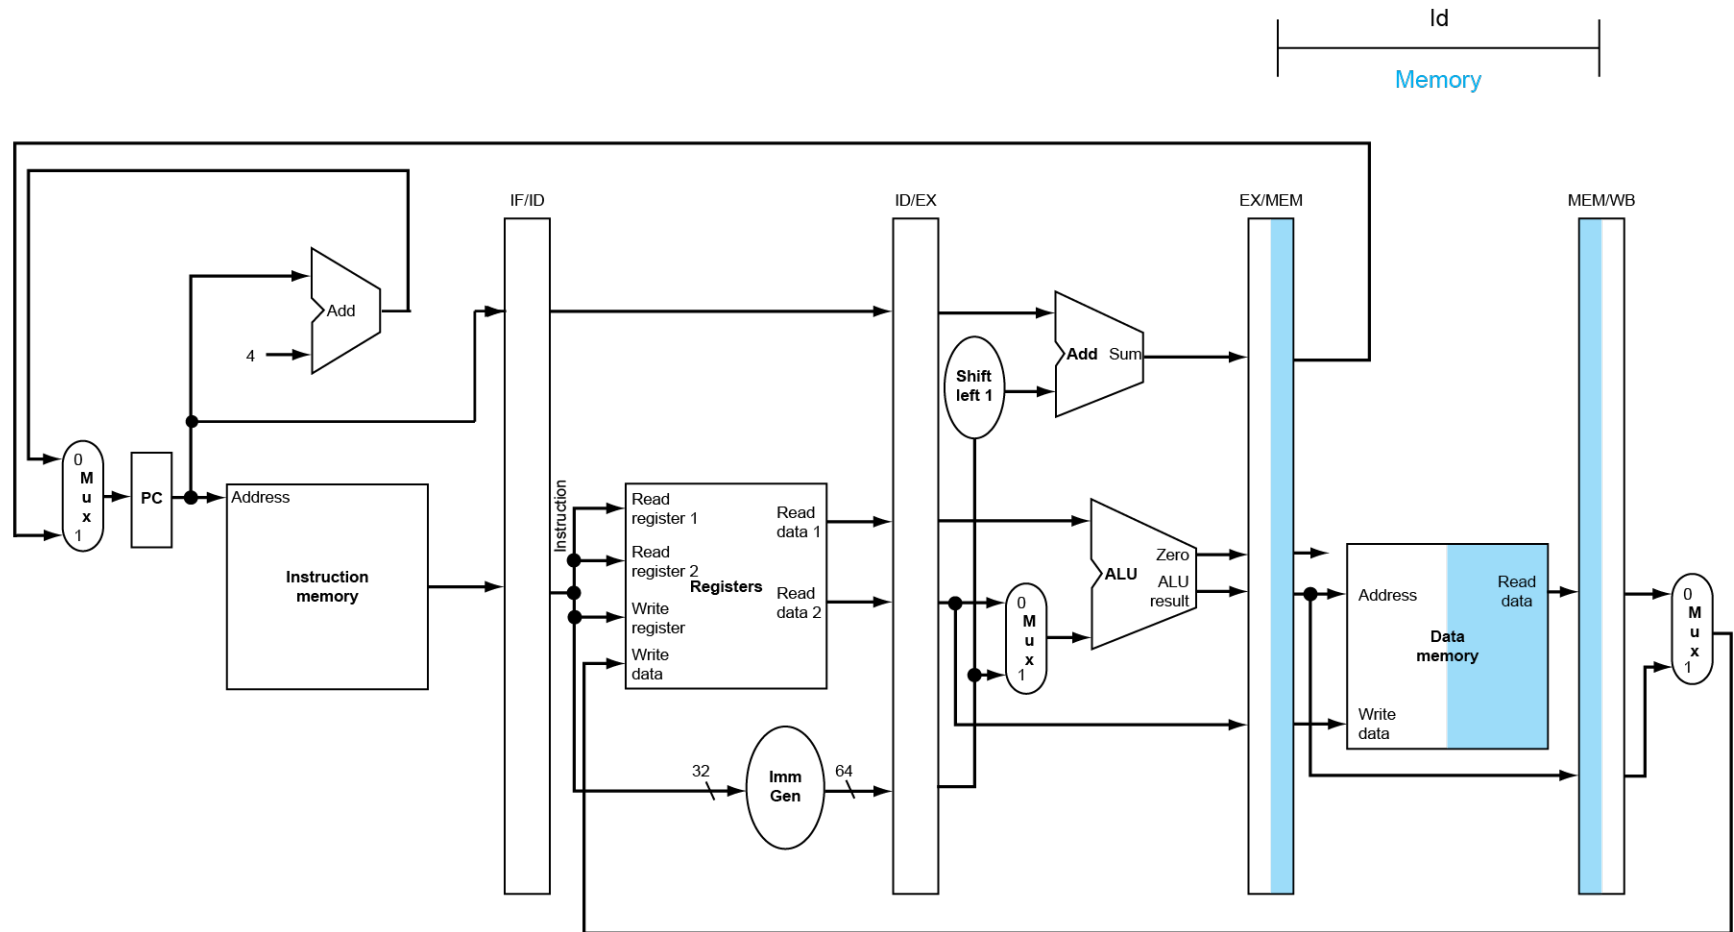
ID for Load, Store, ...



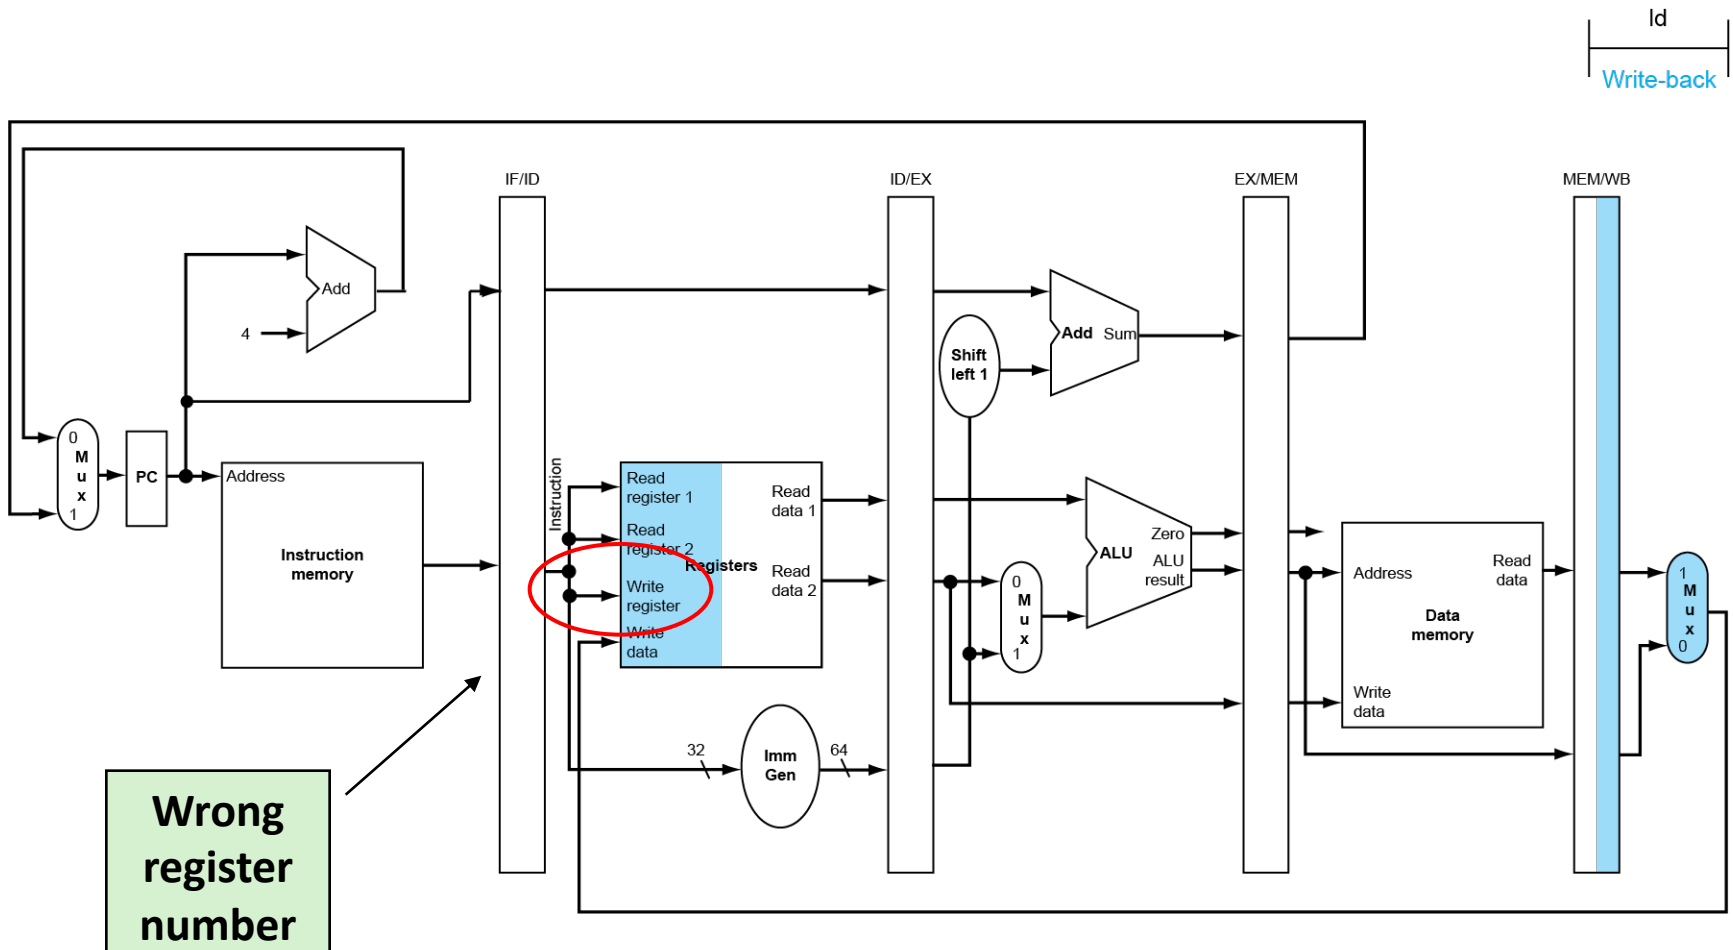
EX for Load



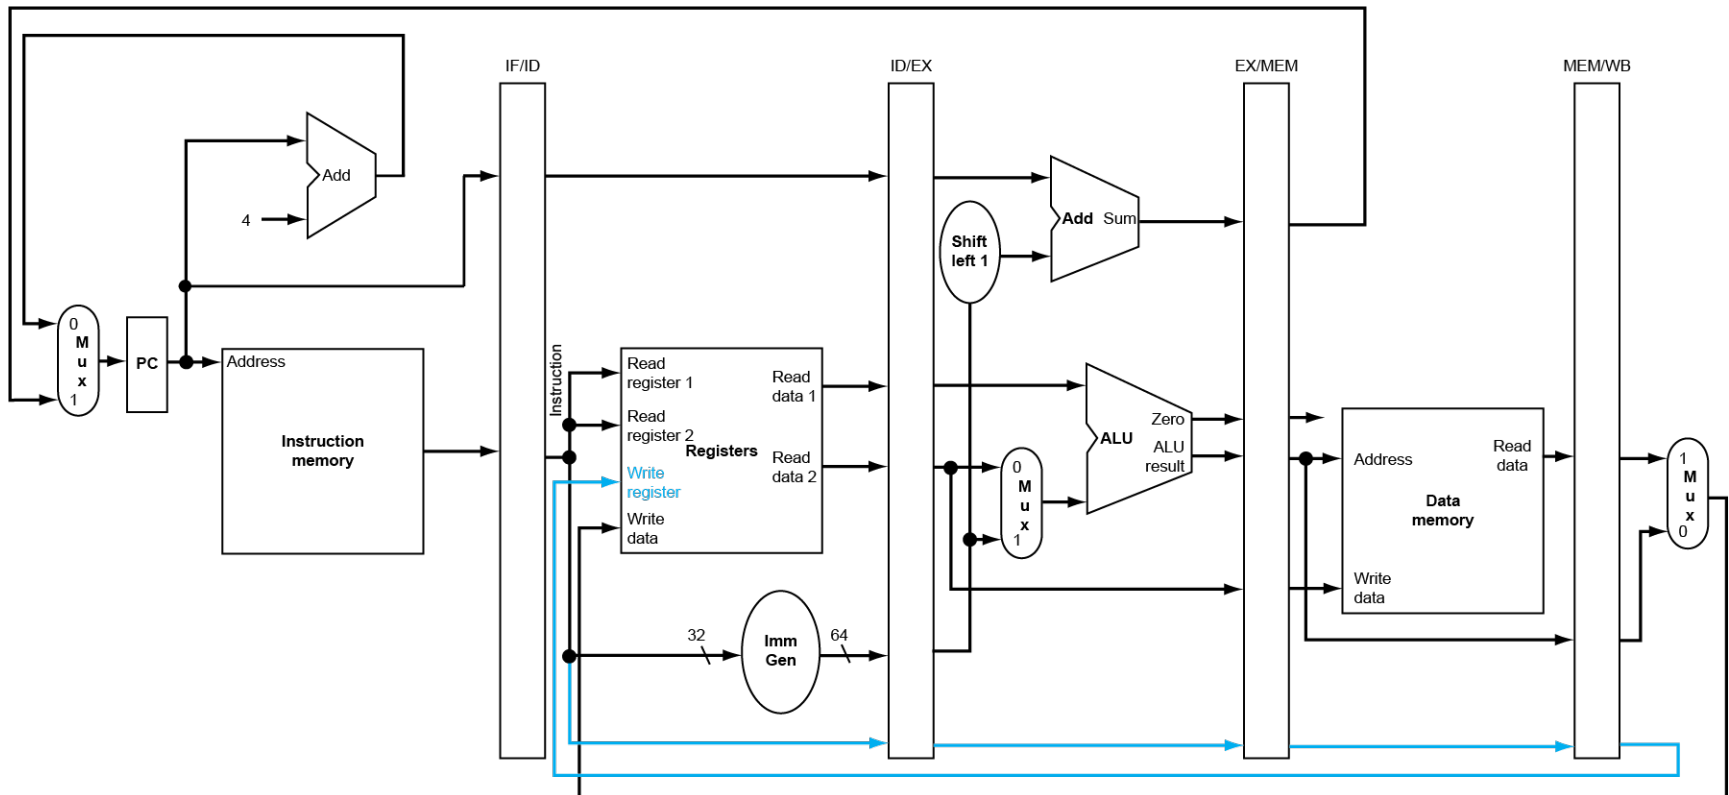
MEM for Load



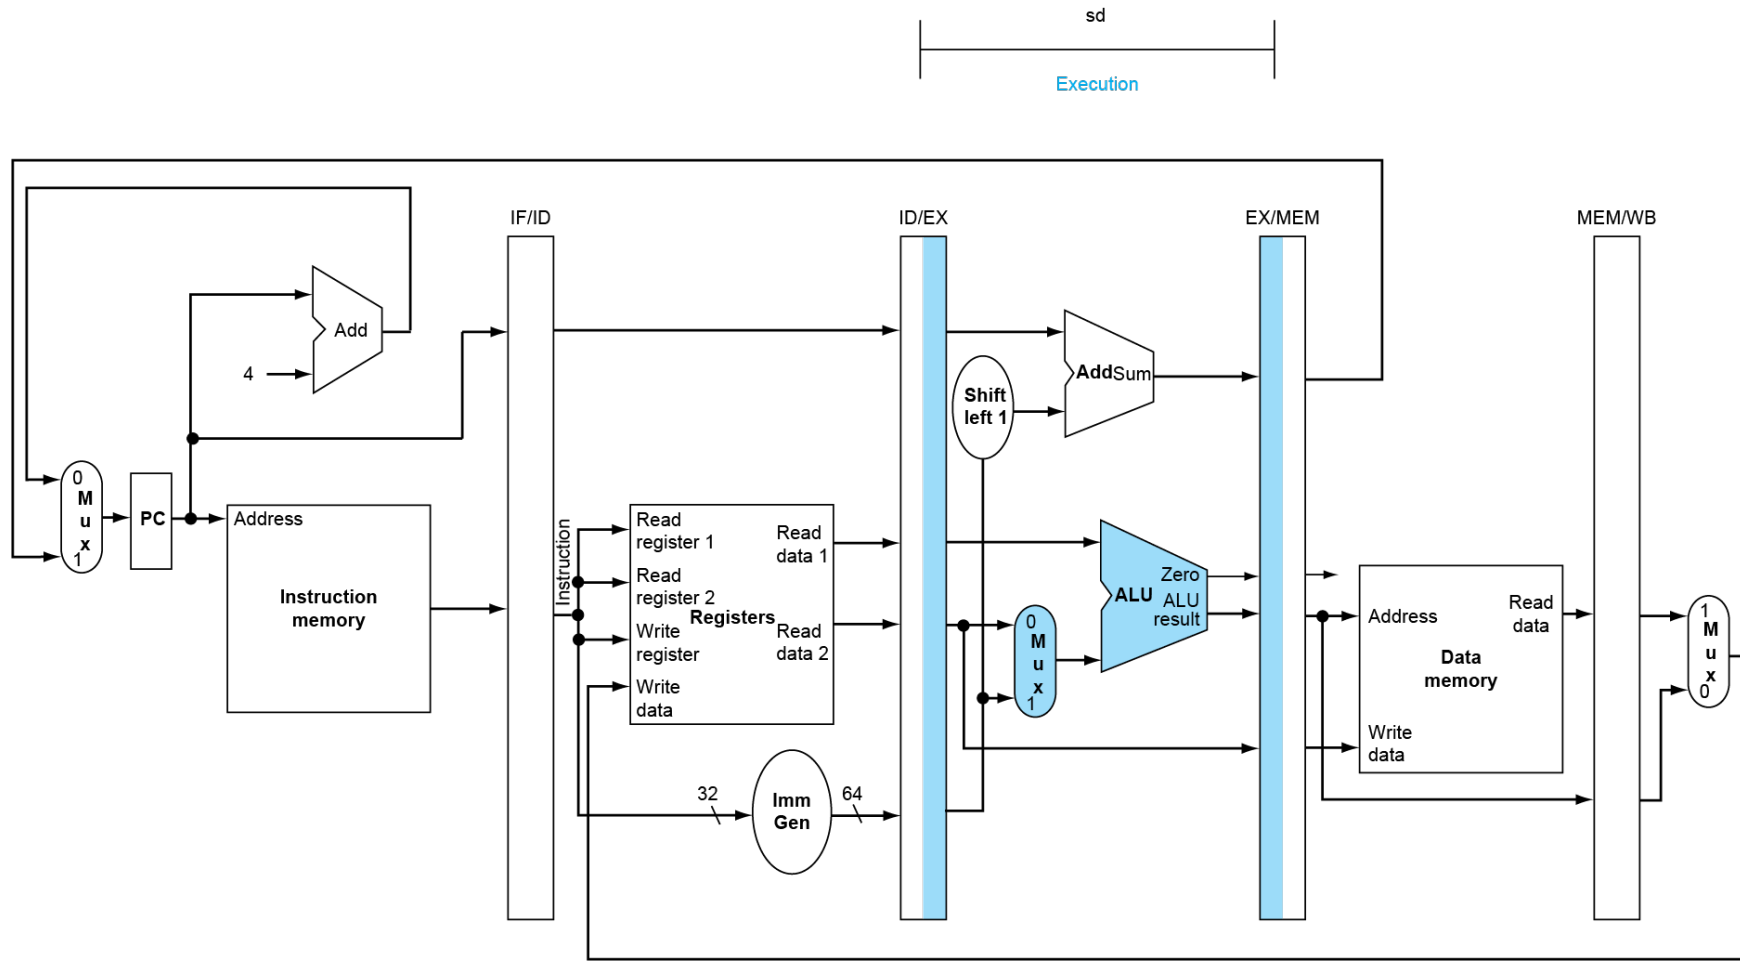
WB for Load



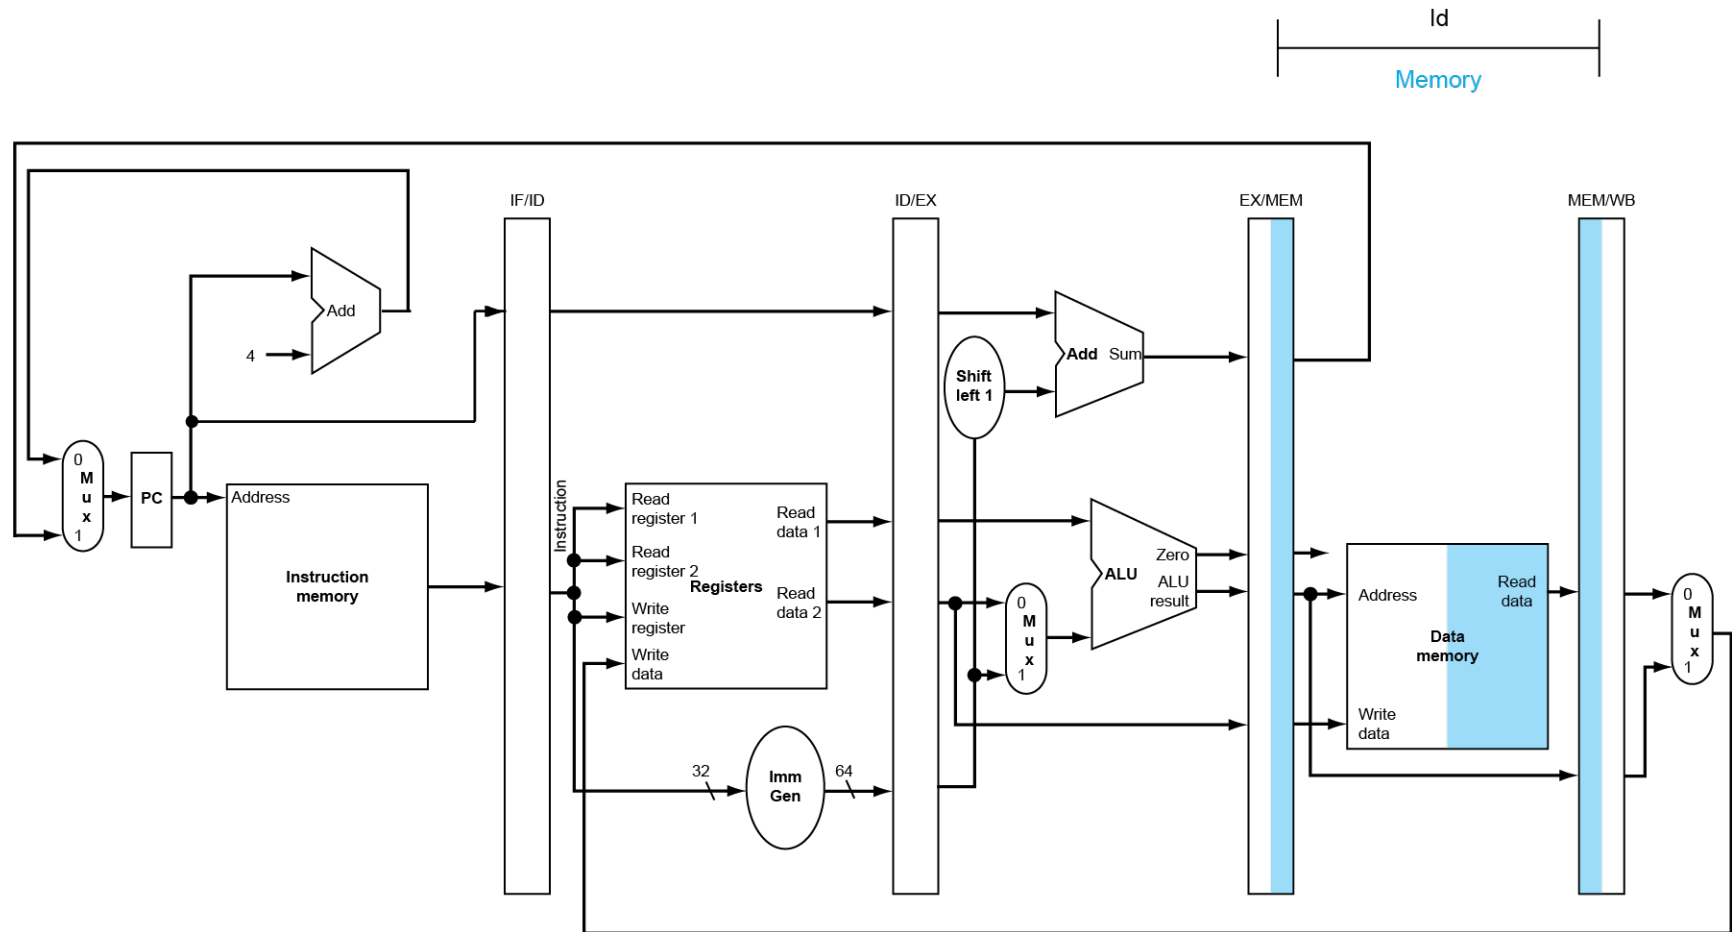
Corrected Datapath for Load



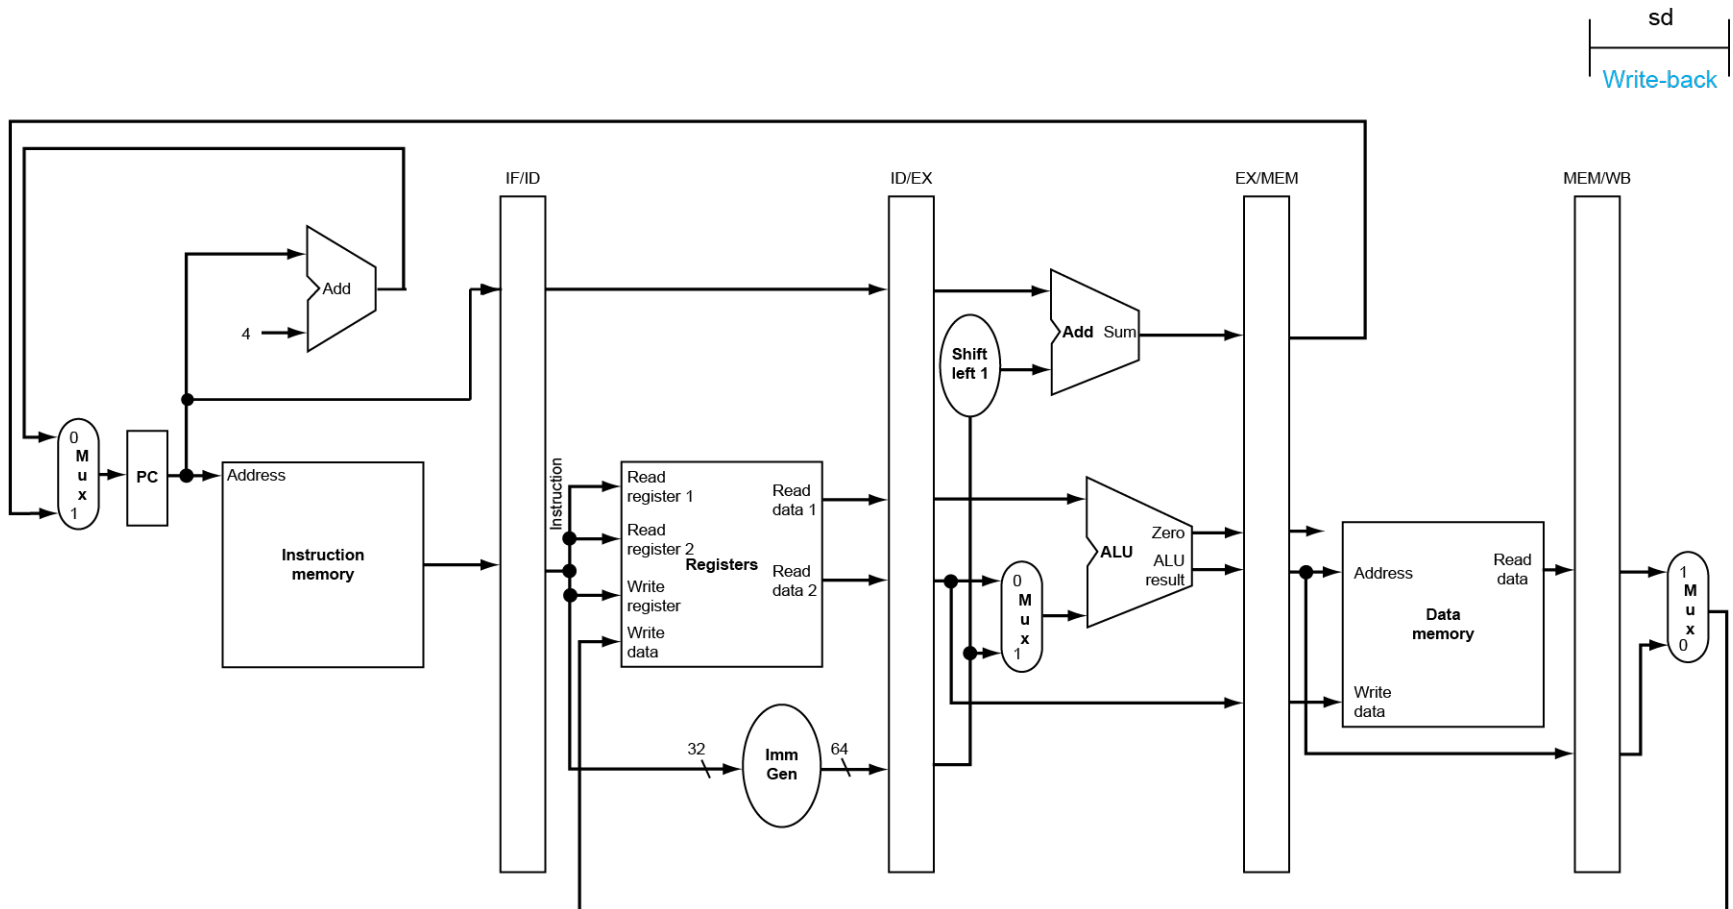
EX for Store



MEM for Store

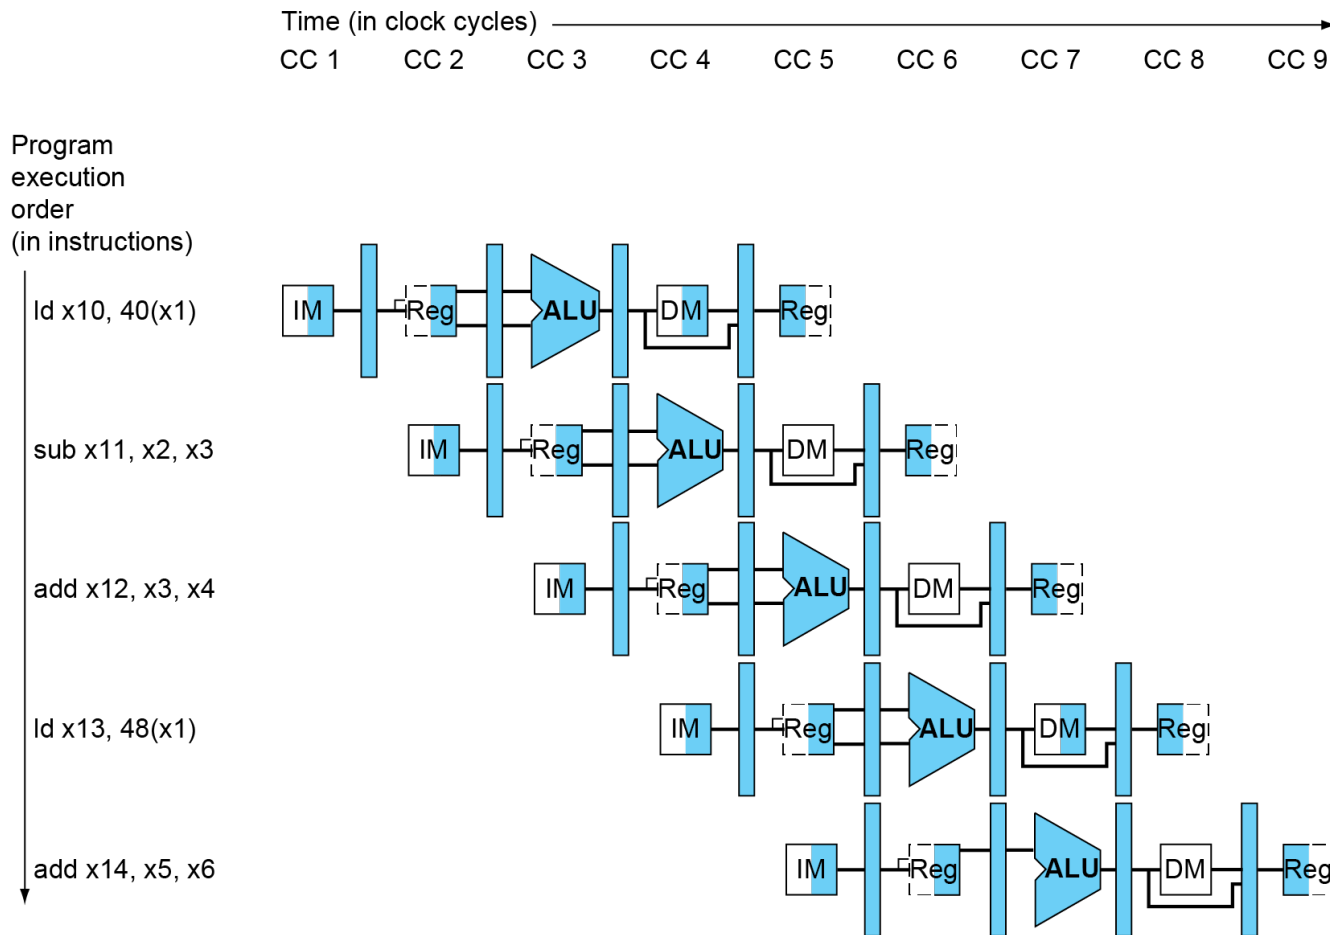


WB for Store



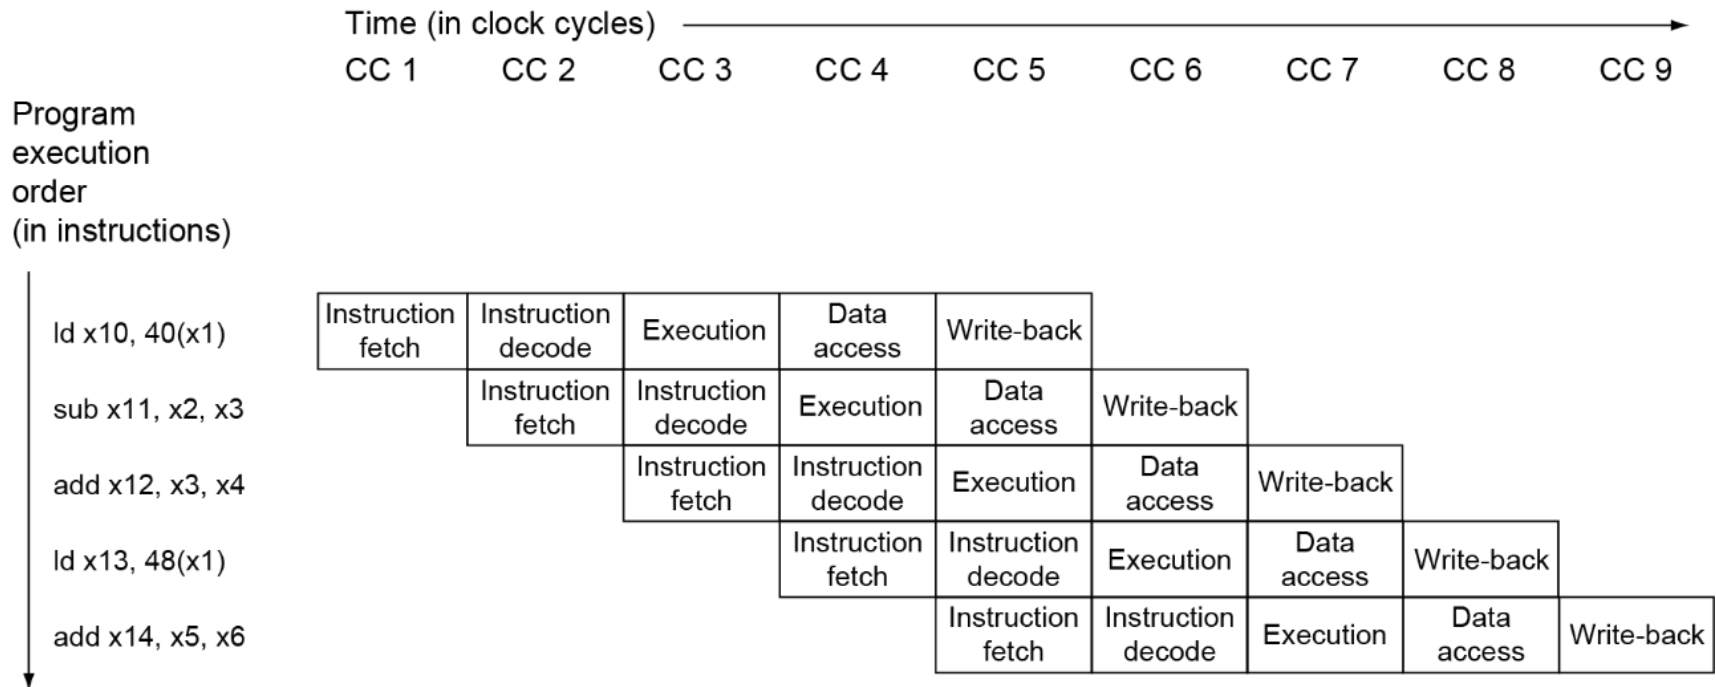
Multi-Cycle Pipeline Diagram

■ Form showing resource usage



Multi-Cycle Pipeline Diagram

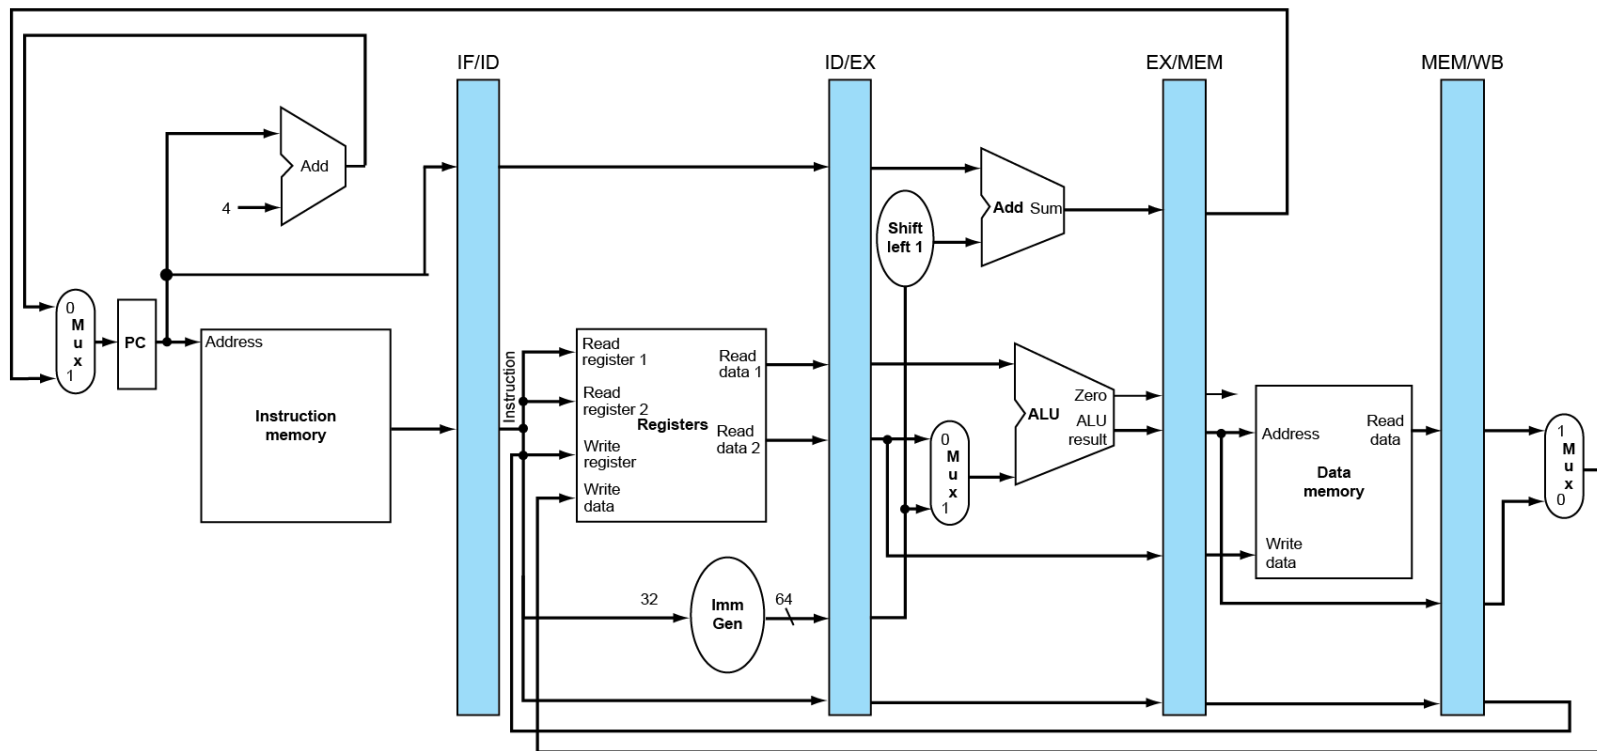
■ Traditional form



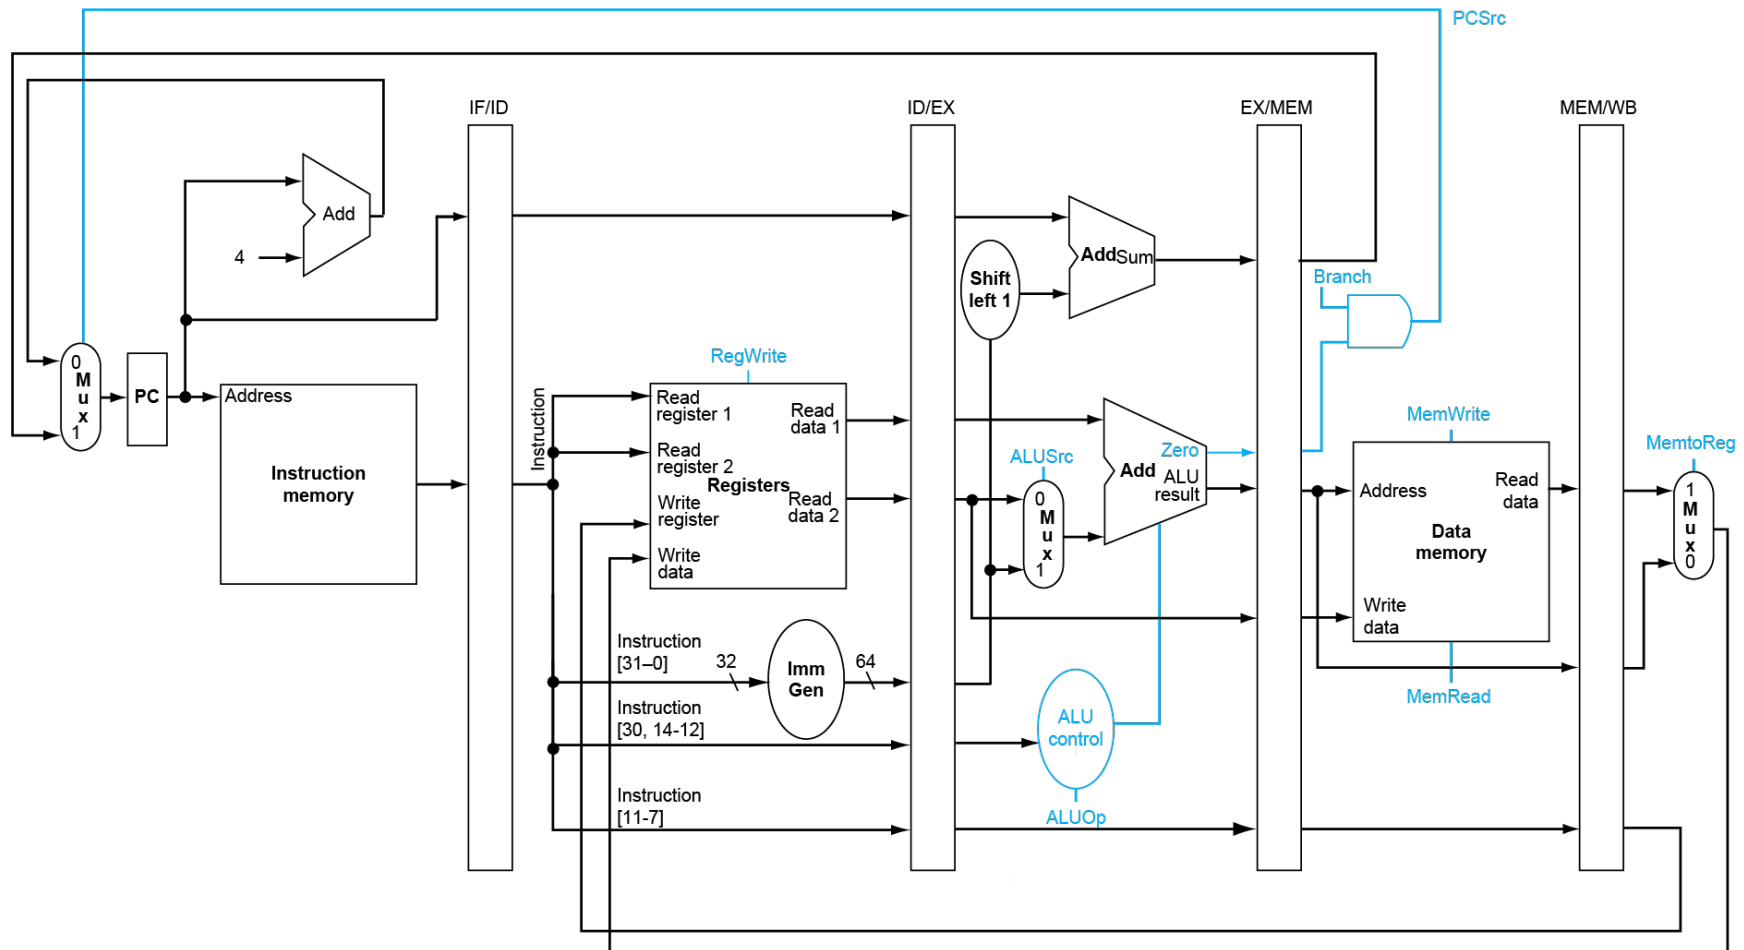
Single-Cycle Pipeline Diagram

■ State of pipeline in a given cycle

add x14, x5, x6	ld x13, 48(x1)	add x12, x3, x4	sub x11, x2, x3	ld x10, 40(x1)
Instruction fetch	Instruction decode	Execution	Memory	Write-back




Pipelined Control (Simplified)



Pipelined Control

- **Control signals derived from instruction**
 - As in single-cycle implementation



Pipelined Control

