

4190.308: Computer Architecture
Midterm Exam
November 1st, 2018
Professor Jae W. Lee

Student ID #: _____

Name: _____

This is a closed book, closed notes exam.

80 Minutes

14 Pages

(+ 2 Pages for Appendices)

Total Score: 200 points

Notes:

- Please turn off all of your electronic devices (phones, tablets, notebooks, netbooks, and so on). A clock is available on the lecture screen.
- Please stay in the classroom until the end of the examination.
- You must not discuss the exam's contents with other students during the exam.
- You must not use any notes on papers, electronic devices, desks, or part of your body.

(This page is intentionally left blank. Feel free to use as you like.)

Part A: Short Answers (16 points)

Question 1 (16 points)

Please answer the following questions. You don't have to justify your answer—just write down your answer only.

Don't guess! You will get 4 points for each correct answer and lose 4 points for each wrong answer (but 0 point for no answer).

- (1) Today's CPUs primarily focus on reducing clock cycle time to improve performance. (True/False)

ANSWER:

- (2) RISC architectures (e.g., MIPS, ARM) generally have an advantage in clock frequency over CISC architectures (e.g., x86-64). (True/False)

ANSWER:

- (3) Unlike integers, the difference between a pair of two adjacent floating-point numbers is non-uniform. (True/False)

ANSWER:

- (4) It is required for a *callee* function to restore the original values of *caller-saved* registers before return to the caller function. (True/False)

ANSWER:

Part B: Integer and Floating-point C Puzzles (24 points)

Question 2 (24 points)

We generate random values for x, y, and z, and convert them to other types on x86-64/Linux/

```
/* Generate random values */
int x = random(); // 4-byte signed
int y = random();
int z = random();

/* Cast to other types */
unsigned ux = (unsigned) x; // 4-byte unsigned
unsigned uy = (unsigned) y;

double dx = (double) x; // double-precision floating point
double dy = (double) y;
double dz = (double) z;
```

For each of the following expressions mark True if it *always* holds; if not, mark False. For every correct answer, you will get 3 points; for every wrong one, you will lose 3 points (and 0 points for no answer).

Expression	True or False?	
$dx * dx \geq 0$	T	F
$(x < y) == (-x > -y)$	T	F
$((x+y) \ll 4) + y - x == 17 * y + 15 * x$	T	F
$ux - uy == -(y - x)$	T	F
$(double)(float)x == dx$	T	F
$dx + dy == (double)(x + y)$	T	F
$dx + dy + dz == dz + dy + dx$	T	F
$dx * dy * dz == dz * dy * dx$	T	F

Part C: Floating-Point Numbers (24 points)

Question 3 (24 points)

To accelerate deep learning applications, Google's hardware team has introduced a new floating-point format, called *bfloat16*, which is a 16-bit floating-point representation based on the IEEE 754 standard. The most significant bit represents a sign bit. The next eight bits are the exponent with a bias of 127. The last seven bits are the fraction. The same rules of the IEEE standard apply (normalized, denormalized, representation of zero, infinity, and NaN).

Sign (1 bit)	Exponent (8 bits)	Fraction (7bit)
-----------------	----------------------	--------------------

(1) Fill in the empty boxes in the following table.

Number	Decimal Representation	Binary Representation
Negative Zero	-0.0	
121/8	15.125 ₁₀	
Positive Infinity	$+\infty$	
One	1	
The smallest negative number		

(2) What is (a) the largest finite number and (b) maximum denormal number for *bfloat16*?

(3) There is another 16-bit floating-point (FP) format called *half-precision FP*, which uses 1, 5, 10 bits for sign, exponent, and fraction, respectively. What is the benefit of *bfloat16* over *half-precision FP*? Explain briefly in a few sentences.
(Hint: Deep learning applications often use *float* type in C to represent data.)

Part D: Human x86-64 Compiler (24 points)

Question 4 (24 points)

Ben Bitdiddle is writing an assembly code, `fib.s` of the original C code (`fib.c`). His code is currently incomplete as the for loop that computes the n -th Fibonacci number is missing. Fill in the missing loop section in `fib.s`. Assume the following register mapping: `x(%rax)`, `y(%rdx)`, `n(%rsi)`, and `i(%rdi)`. *Note that the answer should correct for all integer value n .*

```
/* fib.c */
int main () {
    int x = 0, y = 1, n = 5, z;
    for (int i = 0; i < n; i++) {
        z = x + y;    x = y;    y = z;
    }

    // remaining part (omitted)
    return 0;
}
```

```
# fib.s: x in %rax, y in %rdx, n in %rsi, and i in %rdi
.main
    pushq %rbp
    movq  %rsp, %rbp # initiate procedure
    movq  $0x0, %rax
    movq  $0x1, %rdx
    movq  $0x5, %rsi
    movq  $0x0, %rdi
    # for loop: calculate n-th Fibonacci number into z
    for:

    # remaining part (omitted)
    ...
    mov   $0x0, %rax
    ret
```

Part E: Human x86-64 De-compiler (42 points)

Question 5 (24 points)

The following assembly code shows the body of function `foobar()`.

```

foobar:
    xorq  %r9, %r9          # Initialize i
loop_i:
    movq  %rdi, %rax        # %rdi == arr, %rsi == n
    movq  %r9, %rdx
loop_j:
    movq  (%rdi), %rcx
    movq  (%rax), %r8
    cmpq  %r8, %rcx
    jle   skip
    movq  %r8, (%rdi)
    movq  %rcx, (%rax)
skip:
    addq  $0x1, %rdx
    addq  $0x4, %rax
    cmpq  %rdx, %rsi
    jg    loop_j
    addq  $0x1, %r9
    addq  $0x4, %rdi
    cmpq  %r9, %rsi
    jne   loop_i
    retq

```

Alice Hacker has reconstructed C code from it. Fill in the blanks in the C code below.

```

void foobar(int * arr, int n)
{
    for (int i = 0; i < _____; i++) {
        for (int j = _____; j < _____; j++) {
            if (_____ > _____) {
                int temp = _____;
                _____ = _____;
                _____ = _____;
            }
        }
    }
}

```


Part E: Procedure Calls (30 points)

Question 7 (30 points)

Here is a C program that counts the number of 1's in argument *n*. The assembly code of `popcount()` generated by x86-64/Linux gcc is shown in the right.

Answer the following questions.

<pre> #include <stdio.h> int popcount(int n){ if(n!=0) return (n&0x1) + popcount(n>>1); else return 0; } int main(){ unsigned int n; printf("n : "); scanf("%d",&n); printf("Number of 1's: %d\n", popcount(n)); ② return 0; } </pre>	<pre> 00000000004005f6 <popcount>: 4005f6: push %rbp 4005f7: mov %rsp,%rbp 4005fa: push %rbx 4005fb: sub \$0x18,%rsp 4005ff: mov %edi,-0x14(%rbp) 400602: cmpl \$0x0,-0x14(%rbp) 400606: je 400620 <popcount+0x2a> 400608: mov -0x14(%rbp),%eax 40060b: and \$0x1,%eax 40060e: mov %eax,%ebx 400610: mov -0x14(%rbp),%eax 400613: sar %eax 400615: mov %eax,%edi 400617: callq 4005f6 <popcount> ① 40061c: add %ebx,%eax 40061e: jmp 400625 <popcount+0x2f> 400620: mov \$0x0,%eax 400625: add \$0x18,%rsp 400629: pop %rbx 40062a: pop %rbp 40062b: retq </pre>
--	---

(1) What is the total number of instructions executed if $n=12$ (0x1100)?

- (2) Assuming $n=7$, what are the values of `%ebx`, `%eax`, and `%rip` when ① is reached for the first time?

`%ebx` =

`%eax` =

`%rip` =

- (3) Assuming $n=7$, what will the stack snapshot look like when ① is reached *for the second time*? Fill in the empty table below. Use “???” for an unknown value.

Hints:

- A. `%rsp` and `%rbp` hold `0x7fffffff120` and `0x7fffffff140`, respectively.
- B. The return address to `main()` is `0x400672` (i.e. after all `popcount()` is done).
- C. Just before `popcount(7)` is executed, `%rbp` holds `0x00007fffffff190`.

Stack Address	Value	
	Bytes 7~4	Bytes 3~0
<code>0x7fffffff178</code>	<code>0x00000000</code>	<code>0x00400672</code>
<code>0x7fffffff170</code>		
<code>0x7fffffff168</code>	<code>0x00000000</code>	<code>0x00000000</code>
<code>0x7fffffff160</code>		
<code>0x7fffffff158</code>		
<code>0x7fffffff150</code>		
<code>0x7fffffff148</code>		
<code>0x7fffffff140</code>		
<code>0x7fffffff138</code>	<code>0x00000000</code>	<code>0x00000001</code>
<code>0x7fffffff130</code>		
<code>0x7fffffff128</code>		
<code>0x7fffffff120</code>	<code>0x00000000</code>	<code>0x00000000</code>

Part F: Y86-64 SEQ implementation (40 points)

Question 8 (20 points)

- (1) Using Y86-64 instruction encoding (in Appendix), fill in the box below.
(Note : You may or may not need all 10 bytes(boxes))

Byte 0 1 2 3 4 5 6 7 8 9

(a)

50	15	f8	ff	ff	ff	ff	ff	ff	ff
----	----	----	----	----	----	----	----	----	----

 →

Disassemble

byte 0 1 2 3 4 5 6 7 8 9

(b)

cmovl %rax, %rcx

 →

--	--	--	--	--	--	--	--	--	--

Assemble

- (2) We provide a table of all constants used in HCL below. Complete the HCL code on the next page with correct constants for the Y86-64 SEQ implementation.

Name	Value (Hex)	Meaning
INOP	0	Code for nop instruction
IHALT	1	Code for halt instruction
IRRMovL	2	Code for rrmovl instruction
IIRMOVL	3	Code for irmovl instruction
IRMMOVL	4	Code for rmmovl instruction
IMRMOVL	5	Code for mrmovl instruction
IOPL	6	Code for integer operation instructions
IJXX	7	Code for jump instructions
ICALL	8	Code for call instruction
IRET	9	Code for ret instruction
IPUSHL	A	Code for pushl instruction
IPOPL	B	Code for popl instruction
FNONE	0	Default function code
RESP	4	Register ID for %rsp
RNONE	F	Indicates no register file access
ALUADD	0	Function for addition operation
SAOK	1	Status code for normal operation
SADR	2	Status code for address exception
SINS	3	Status code for illegal instruction exception
SHLT	4	Status code for halt

Decode		<pre> int srcA = [icode in { _____ _____ } : rA; icode in { _____ _____ } : RRSP; 1 : _____; # Don't need register]; </pre>
Execute		<pre> int aluA = [icode in { IRRMOVQ, IOPQ } : valA; icode in { _____ _____ } : valC; icode in { _____ _____ } : -8; icode in { IRET, IPOPQ } : 8; # Other instructions don't need ALU]; int alufun = [icode == _____ : ifun; 1 : _____]; </pre>
Memory		<pre> int Stat = [imem_error dmem_error: _____; !instr_valid: _____; icode == IHALT : _____; 1 : _____;]; bool mem_read = icode in { _____ _____ }; </pre>

Question 9 (20 points)

We would like to add `addmem` instruction to the Y86-64 sequential implementation. `Addmem` instruction takes 2 operands, one from register and the other from the memory, and performs an addition. In other words, `addmem D(rB), rA` computes $rA \leftarrow rA + \text{Mem}[rB + D]$. Answer the following questions.

	icode:fn		rA:rB		D(8byte)
addmem rA,D(rB)	C = ADDMEM	0	rA	rB	D

(1) Describe the additional hardware (datapath) required to implement this instruction in one paragraph.

(2) What is the impact of adding this instruction on the CPU cycle time? Would it be increased or unchanged? Again, explain in one paragraph.

(This page is intentionally left blank. Feel free to use as you like.)

Appendix A: X86-64 assembly

Common instructions

mov	src, dst	dst = src
movsbl	src, dst	byte to int, sign-extend
movzbl	src, dst	byte to int, zero-fill
lea	addr, dst	dst = addr
add	src, dst	dst += src
sub	src, dst	dst -= src
imul	src, dst	dst *= src
neg	dst	dst = -dst (arith inverse)
sal	count, dst	dst <= count
sar	count, dst	dst >= count (arith shift)
shr	count, dst	dst >= count (logical shift)
and	src, dst	dst &= src
or	src, dst	dst = src
xor	src, dst	dst ^= src
not	dst	dst = ~dst (bitwise inverse)
cmp	a, b	b-a, set flags
test	a, b	a&b, set flags
jmp	label	jump to label (unconditional)
je	label	jump equal ZF=1
jne	label	jump not equal ZF=0
js	label	jump negative SF=1
jns	label	jump not negative SF=0
jg	label	jump > (signed) ZF=0 and SF=OF
jge	label	jump >= (signed) SF=OF
jl	label	jump < (signed) SF!=OF
jle	label	jump <= (signed) ZF=1 or SF!=OF
ja	label	jump > (unsigned) CF=0 and ZF=0
jb	label	jump < (unsigned) CF=1
push	src	add to top of stack Mem[--%rsp] = src
pop	dst	remove top from stack dst = Mem[%rsp++]
call	fn	push %rip, jmp to fn
ret		pop %rip

Instruction suffixes

b	byte
w	word (2 bytes)
l	long/doubleword (4 bytes)
q	quadword (8 bytes)


Suffix is elided when can be inferred from operands
e.g. operand %rax implies q, %eax implies l, and so on

IEEE 754 FLOATING-POINT STANDARD

$$(-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

where Single Precision Bias = 127,
Double Precision Bias = 1023.

IEEE Single Precision and Double Precision Formats:



Condition flags

ZF	Zero flag
SF	Sign flag
CF	Carry flag
OF	Overflow flag

④
IEEE 754 Symbols

Exponent	Fraction	Object
0	0	± 0
0	≠ 0	± Denorm
1 to MAX - 1	anything	± FL Pt. Num.
MAX	0	±∞
MAX	≠ 0	NaN

S.P. MAX = 255, D.P. MAX = 2047

Registers

%rip	Instruction pointer
%rsp	Stack pointer
%rax	Return value
%rdi	1st argument
%rsi	2nd argument
%rdx	3rd argument
%rcx	4th argument
%r8	5th argument
%r9	6th argument
%r10, %r11	Caller-saved
%rbx, %rbp, %r12...%r15	Callee-saved

Addressing modes

Example source operands to **mov**

Immediate

mov \$0x5, dst

\$val

source is constant value

Register

mov %rax, dst

%R

R is register

source in %R register

Direct

mov 0x4033d0, dst

0xaddr

source read from Mem[0xaddr]

Indirect

mov (%rax), dst

(%R)

R is register

source read from Mem[%R]

Indirect displacement

mov 8(%rax), dst

D(%R)

R is register

D is displacement

source read from Mem[%R + D]

Indirect scaled-index

mov 8(%rsp, %rcx, 4), dst

D(%RB, %RI, S)

RB is register for base

RI is register for index (0 if empty)

D is displacement (0 if empty)

S is scale 1, 2, 4 or 8 (1 if empty)

source read from

Mem[%RB + D + S*%RI]

** Originally from Stanford CS107;
modified for SNU CSE 4190.308*

Appendix B: Y86-64 (Instruction Set)

Instruction	icode:fn		rA:rB						
	byte	0	1	2	3	4	5	6	7 8 9
halt		0 = IHALT	0						
nop		1 = INOP	0						
cmovXX rA, rB		2 = IRRMOVQ	fn						
rrmovq			0						
cmovle			1						
cmovl			2						
cmove			3						
cmovne			4						
cmovge			5						
cmovg			6						
irmovq V, rB		3 = IIRMOVQ	0	F	rB	V			9
rmmovq rA, D(rB)		4 = IRMMOVQ	0	rA	rB	D			
rrmovq D(rB), rA		5 = IMRMOVQ	0	rA	rB	D			
OPq rA, rB		6 = IOPQ	fn	rA	rB				
addq			0						
subq			1						
andq			2						
xorq			3						
jXX Dest		7 = IJXX	fn	Dest					8
jmp			0						
jle			1						
jl			2						
je			3						
jne			4						
jge			5						
jg			6						
call Dest		8 = ICALL	0	Dest					8
ret		9 = IRET	0						
pushq rA		A = IPUSHQ	0	rA	F				
popq rA		B = IPOPQ	0	rA	F				

Register encoding

0	1	2	3	4	5	6	7
%rax	%rcx	%rdx	%rbx	%rsp	%rbp	%rsi	%rdi
8	9	A	B	C	D	E	F
%r8	%r9	%r10	%r11	%r12	%r13	%r14	No register